

How Spillovers Erode Participation in Conservation Buybacks: Evidence from the Texas Inshore Shrimp Fishery

Eric Belleville
Ohio State University
Belleville.11@buckeyemail.osu.edu

September 22, 2026

Abstract

Buyback programs are widely used to reduce overcapitalization and generate social benefits in settings such as land retirement, groundwater conservation, and fisheries. Yet evaluations often stop at program boundaries, even though removals can generate spillovers: displaced activity may move elsewhere, while those who remain benefit from a replenished resource. I study how these gains feed back into buyback program participation. As capacity is removed, the continuation value of remaining participants rises, creating an incentive to hold out rather than exit. Successful removals can therefore weaken the program participation required for further success.

I quantify this feedback-loop using the Texas Parks and Wildlife Department's 26-year inshore shrimp license buyback as a testbed. I estimate a two-stage structural model of vessels' spatial fishing choices and license-retirement decisions, then use it to simulate continuation values and reconstruct harvest and auction activity when vessel-level data are unavailable. I find that for every thirteen active license-holders who bid, roughly one who would otherwise have participated holds out as stock recovery makes staying more valuable. Program effectiveness is further constrained by two margins: approximately two-thirds of retired licenses were already inactive, and surviving vessels replace more of the retired effort in denser bay systems. At the same budget, restricting purchases to active licenses would have retired 123 more trawl-days per license, while targeting San Antonio bay alone would have retired 306 more. Together, the results show that conservation buybacks can undermine participation and additionality through the benefits they create for those who remain.

Contents

1	Introduction	1
2	Institutional Background	6
2.1	The Texas Inshore Shrimp Fishery	6
2.2	The Texas Shrimp License Buyback Program	6
3	Data	9
4	A Dynamic Model of Common-Pool Exploitation	11
4.1	Model Overview	11
4.2	Production Function and Stock Dynamics	14
4.2.1	Harvest Production	14
4.2.2	Latent Stock Index	15
4.2.3	Within-Season Growth Function	15
4.3	Bay System Choice	16
4.4	State Variable Transitions	16
4.5	Equilibrium and Continuation Values	17
4.5.1	Markov Perfect Equilibrium	17
4.5.2	Forward Simulation	18
4.5.3	Continuation Value Computation	19
5	Estimation Results: Fishery Dynamics	20
5.1	Production Function Estimates	20
5.2	Stock Dynamics	20
5.3	State Transition Fit	22
5.4	Stage 1 Conditional Choice Probabilities	23
5.5	Structural Profit Parameters	24
5.6	Continuation Values and Auction Entry	27
6	Counterfactual Analysis: Spillovers, Offsets, and Targeting	28
6.1	Overview	28
6.2	The Synthetic Fleet and Full-Horizon Back-Cast	29
6.3	The Participation Distortion and Its Decomposition	32
6.4	The Structural Offset: How Much Conservation a Removal Delivers	34
6.5	The Cost-Effectiveness of Removal Rules	36
7	Conclusion	37
A	Supplementary Estimation Results	44
A.1	Vessel Atomicity	44
A.2	Stock-Index Components	44
A.3	Bay-Shrimp Inverse Demand	44
A.4	Auction Entry: Risk-Set Robustness	44

B Synthetic-Fleet Construction and Full-Horizon Back-Cast	47
C Counterfactual Supplementary Exhibits	49
D Model Validation	51

1 Introduction

Buyback programs—reverse auctions and fixed-price offers that pay resource users to retire environmentally harmful capital—have become a mainstream conservation instrument. Programs in the United States alone have spent billions each year on cropland retirements [U.S. Department of Agriculture, Farm Service Agency, 2025a], and millions on several fishing license buyback programs [NOAA Fisheries, 2024a] and groundwater rights retirements [Sakas, 2022]. Their appeal over other instruments such as taxes and quotas is that they are voluntary: participants reveal, and are compensated for, their privately known cost of exit, relieving an overcapitalized or economically fragile resource without forcing out vulnerable users [Ferraro and Kiss, 2002, Jack et al., 2008]. But when participation falls short, the social benefits fall with it, delivering far less per dollar when participant pools thin [DePiper, 2015, Rolfe et al., 2018]. When the Conservation Reserve Program (CRP) falls 1 million acres short of its statutory cap [U.S. Department of Agriculture, Farm Service Agency, 2025a], or recruits’ retirements that are only 25% additional [Aspelund and Russo, 2026], it may be read as a failure of price, outreach, or targeting. But what if withholding participation is a rational response to the program’s own success?

A large literature accounts for how the effects of conservation programs spill beyond their targets, often called “slippage” or “leakage” effects. For example, when a nation commits to a carbon abatement goal, non-committed countries gain a competitive advantage in carbon-intensive industries [Aichele and Felbermayr, 2015, King and van den Bergh, 2021]; and when a landowner sells cropland for non-farm uses, they are likely to bring other idle land back into farming [Wu, 2000]. Pfaff and Robalino [2017] organize these mechanisms into several channels, including input reallocation, market prices, learning, and non-pecuniary, which each elicit a behavioral response. A fifth channel, “ecological-physical links”, is largely set aside because it “... does not involve human behavior”, though the authors note that “... those links may also lead to further behavioral response and thus additional program spillovers.” Examples include ecosystem-service gains to neighboring farms when cropland is retired to grassland [Landis et al., 2008], increasing pumping by connected groundwater users after nearby rights are retired [Rouhi Rad et al., 2021, Ekpe and Klis, 2023], and reallocated fishing activity after quota programs restrict harvest in a primary fishery [Cunningham et al., 2016] (See Table 1 for more examples). This paper focuses on this fifth channel, but on a margin that has received little attention: ecological recovery changes the private continuation value of agents who remain eligible for the program. I ask how that endogenous gain feeds back into participation, and how program design can account for it.

While the programs above generate positive externalities for those who remain, auction theory, which serves as the conceptual basis for many analyses in this space, has more often studied strategic distortions arising from negative externalities. When a rival’s win is costly, bidders compete more aggressively, which can increase the principal’s revenue [Hu et al., 2013]. This logic appears in settings such as timber rights [Kuehn, 2019], research patents [Kwon, 2020], and security auctions [Hernandez-Chanto and Fioriti, 2019], where winning confers an advantage in downstream competition. In a conservation buyback, the sign of the externality reverses: when a rival’s exit improves conditions for those who remain, each participant would prefer that someone else sell instead. Positive externalities can therefore

weaken participation and raise the cost¹ of inducing further exit [Calel, 2012, Espinola-Arredondo, 2008, Choi et al., 2018]. The same holdout logic appears in models of voluntary environmental agreements, where the incentive to remain outside a coalition grows as more countries commit to abatement [Barrett, 1994, Nordhaus, 2015].

Because voluntary procurement compensates agents for their private cost of exit, and that cost is lowest for those who would have scaled back regardless, a parallel concern emerges when a program’s budget pays for conservation that would have happened without it. This adverse selection is the central efficiency problem in conservation contracting: it has been studied as contract design under asymmetric information [Ferraro, 2008], screening when additionality is unobserved [Mason and Plantinga, 2013], and measurement against a counterfactual of rejected bids [Rosenberg and Pratt, 2024, Aspelund and Russo, 2026]. This work treats additionality as a property of the individual contract, but when benefits spillover on agents who remain the addition of a retirement becomes endogenous. But in a common-pool resource, the conservation delivered is attenuated further by how much of the resource the remaining users recapture, so even a removal that is fully additional in the usual sense may retire less than it appears.

This recapture, the offset, is itself a familiar object. It has been measured as reduced-form leakage in fisheries [Cunningham et al., 2016] and modeled as the redistribution of effort across patches in spatial bioeconomics [Sanchirico and Wilen, 2020], and the buyback policy literature has long warned that these programs are blunted by the anticipation of future removals [Clark et al., 2005], by the replacement of retired capacity [Holland et al., 1999], or by fishermen overinvesting in gear in anticipation of a thinned rival pool, known as “capital stuffing” [Clark and Munro, 2002]. The same recapture exists wherever the rights are spatial and the resource is shared, from effort displacement in rights-based fisheries [Abbott and Wilen, 2009] to pumping externalities across a common aquifer [Pfeiffer and Lin, 2012].

Several issues make identification in this vein challenging. For example, the debate on the magnitude of physical spillovers in the CRP has persisted for decades, with credible estimates ranging from null to substantial depending on how econometricians choose to handle issues like land heterogeneity, endogenous entry, and unobservable variables [Wu, 2000, Roberts and Bucholtz, 2005, Fleming, 2014]. One particular challenge emerges from the wide breadth of spillover channels outlined by Pfaff and Robalino [2017]: it is challenging to isolate congestion, price, and physical links when entry into programs is endogenously determined by an individual’s continuation value. To bridge this empirical gap, I frame the Texas Parks and Wildlife Department’s (TPWD) inshore shrimp license buyback, a reverse auction run over 41 rounds and 26 years across eight bay systems, as a testbed. This setting provides two unique comparative advantages for identification over previous buyback settings. First, the spillover effect is tractable and structurally identifiable; the removal of licenses leads to a decrease in cumulative harvest, which in turn improves stock health and allows remaining vessels to harvest more shrimp at lower cost. Second, it isolates a biological spillover. Congestion relief could in principle generate its own positive spillover, but the bay systems span thousands of square kilometers and host numerous other fisheries; I show using

¹In a forward auction, bids reflect willingness to pay and the principal seeks high bids; in a reverse auction, asks reflect willingness to sell and the principal seeks low asks. The strategic implications of externalities therefore translate with the direction of the mechanism.

bay-specific shift-share instruments that the congestion effect is negligible and insignificant. Price-mediated leakage is likewise minimal, as Gulf-harvested shrimp make up only 4.5% of the U.S. food-shrimp market [Liese et al., 2026].

First, I provide evidence that a positive stock externality can distort participation within a repeated conservation buyback. Leakage and slippage are familiar in conservation policy: program gains erode as activity shifts to unregulated places, agents, or periods [Wu, 2000, Cunningham et al., 2016]. The literature primarily studies spillovers outward from the program boundary. I instead study a feedback loop within the program itself: by removing effort, the buyback rebuilds the stock and raises the continuation value of vessels that remain, making future exit less attractive. The TPWD setting allows me to isolate this biological channel from price and congestion effects. A one-million-pound increase in cumulative seasonal harvest lowers stock health by 47%, and a higher continuation value reduces bidding propensity by about 4%; over the full program, stock recovery suppresses participation by roughly 7.6%.

Second, I extend the dynamic discrete-choice framework for common-pool exploitation of Aguirregabiria and Mira [2002, 2007], brought to fisheries by Huang and Smith [2014], to a spatial, multi-bay setting, and use the estimated model to reconstruct participation and harvest over the program’s full twenty-six-year horizon, including years for which vessel-level activity data is unavailable. Third, I unify additionality and offset into a single measure of conservation delivered per dollar, and derive the offset margin structurally from the resource externality itself. Asset retirement falls short on two margins. The first is extensive and familiar from land-retirement programs: a purchased license may have been inactive to begin with, so its retirement is non-additional [Ferraro, 2008, Jack et al., 2008]. The second is intensive and specific to a common-pool resource: even a fully active retirement delivers less than one-for-one conservation because surviving vessels fish the freed stock and recapture part of the removed effort. I recover this offset directly from the estimated depletion technology and show that it is not constant, but rises with local fleet density from 0.04 in the thinnest bay to 0.26 in the densest as recapture saturates with remaining activity. Together, these margins imply large heterogeneity in effective additionality and conservation per dollar: screening purchases for recent activity and targeting low-offset bays substantially increase net trawl-days retired for the same budget.

The paper proceeds as follows. Section 2 describes the Texas inshore shrimp fishery and the buyback program. Section 3 details the data. Section 4 develops the two-stage structural model with harvest production, latent stock dynamics, state-variable transitions, and the equilibrium used to simulate continuation values. Section 5 reports estimates and traces the spillover-to-participation chain. Section 6 simulates counterfactual program designs and isolates the participation distortion from retirement-driven stock replenishment; Section 7 concludes, and supporting tables and figures are listed in Section A.

Table 1: Buyback Settings, Spillovers, and Feedback Loops

Setting (retired unit)	Spillover	Recapture and participation feedback	Flagship voluntary program (approx. scale)
Marine & coastal fisheries (licenses, vessels)	Removal slows depletion, raising the stock. <i>This paper</i> : depletion semi-elasticity ≈ 0.64	Remaining vessels intensify, rising with density (0.01–0.52).	NOAA capacity buybacks: Pacific groundfish $\approx \$46\text{M}$, AFA pollock $\approx \$75\text{M}$ [NOAA Fisheries, 2024b]; TX inshore shrimp $\approx \$15\text{M}$ (<i>this paper</i>)
Groundwater (acreage, pumping rights)	Less pumping raises the water table, cutting cost and raising yield for connected users	Neighbors pump more [Rouhi Rad et al., 2021].	CREP (Republican River); SGMA following $\approx 0.5\text{--}0.9\text{M}$ ac [Public Policy Institute of California, 2023]; CO buy-outs $\approx \$60\text{M}$ [Sakas, 2022]
Surface water / river flow (entitlements, diversions)	Retired diversions leave more in-stream and downstream flow	Remaining and junior users take the freed water.	Murray–Darling Basin water buybacks $\approx \text{A\$}2.5\text{B}$, $\sim 1,200$ GL [Australian Government Department of Climate Change, Energy, the Environment and Water, 2024]
Rangeland / grazing (permits, AUMs)	Less grazing improves forage and riparian condition for remaining permittees	Remaining herds expand.	Voluntary Grazing Permit Retirement (US; mostly proposed) [U.S. House of Representatives, 2025]
Cropland: pollination & pest control (cropland to habitat)	Habitat raises pollination and natural pest control on adjacent farms (local, diffuse)	Slippage onto the owner’s other land [Wu, 2000].	CRP, incl. pollinator practice CP42, $\approx \$1.6\text{--}1.8\text{B/yr}$, $\sim 23\text{M}$ ac [Aspelund and Russo, 2026]
Cropland / watershed: water quality (cropland, buffers)	Less runoff yields cleaner water and less sediment downstream	Slippage.	CREP riparian buffers; Chesapeake Bay nutrient trading [U.S. Department of Agriculture, Farm Service Agency, 2025b]
Wetlands (cropland to restored wetland)	Flood attenuation, water quality, and habitat for neighbors	Limited.	Wetland Reserve / ACEP-WRE, $\sim 2\text{--}3\text{M}$ acres [U.S. Department of Agriculture, Natural Resources Conservation Service, 2025]
Carbon / emissions (emissions, allowances, credits)	Avoided emissions, a non-rival climate benefit enjoyed regardless of participation	Geographic leakage [Aichele and Felbermayr, 2015, King and van den Bergh, 2021]. Feedback: absent (competitiveness channel separate)	EU ETS; RGGI; California cap-and-trade; CDM / REDD+ and voluntary off-sets

Notes: Acronyms, in alphabetical order. ACEP-WRE: Agricultural Conservation Easement Program, Wetland Reserve Easements (USDA). AFA: American Fisheries Act. AUM: animal unit month, the forage required by one cow-calf pair for one month. CDM: Clean Development Mechanism. CP42: CRP conservation practice 42, pollinator habitat. CREP: Conservation Reserve Enhancement Program. CRP: Conservation Reserve Program. EU ETS: European Union Emissions Trading System. GL: gigaliter, one billion liters. NOAA: National Oceanic and Atmospheric Administration. REDD+: Reducing Emissions from Deforestation and Forest Degradation, including conservation, sustainable forest management, and enhancement of forest carbon stocks. RGGI: Regional Greenhouse Gas Initiative. SGMA: Sustainable Groundwater Management Act (California). A\$ denotes Australian dollars; CO and TX denote Colorado and Texas.




Figure 1: Texas Inshore Shrimp Bay Systems. Bay-shrimp landings (lbs) are displayed from 2008 to 2023 ticketed Bay shrimp landings. Five bay systems outlined in red are the focal point for this paper: Galveston, Matagorda, San Antonio, Aransas, and Corpus Christi bay systems (read North to South).

2 Institutional Background

2.1 The Texas Inshore Shrimp Fishery

The Texas inshore shrimp fishery, regulated by the Texas Parks and Wildlife Department (TPWD), is composed of eight major bay systems along the Gulf of Mexico, comprising an area of roughly 4,363 square kilometers (see Figure 1). Roughly 293 million pounds of bay shrimp tails have been landed since 1981, generating a cumulative total revenue of \$1.2 billion (in adjusted 2023 US dollars). While these bay systems are large, operations are still heavily interconnected between vessels. Shrimp populations need specific levels of salinity and narrow corridors to thrive. Biomass stock can also be heavily influenced by active licenses, as less-exploited populations can thrive. While vessels are not bumping into each other, they generally can be assumed to enjoy positive returns from less capitalized bay systems.

Inshore shrimp refers to two broad species categories: “Bay” shrimp, the focus of this paper, is commercially produced for food, while “bait” shrimp is primarily used as an input toward other fisheries. Because many state laws prohibit the use of imported shrimp as bait, Gulf bait shrimp carries huge market power domestically [Texas Parks and Wildlife Department, 2026]. The bay shrimp market, on the other hand, faces substantial pressure from foreign imports such that vessels can safely be assumed to be price takers² [Liese et al., 2026]. TPWD also manages deep-sea shrimp fisheries which they refer to as “Gulf” shrimp, but they have not sought to reduce capitalization in those fisheries. The goal of this paper is to analyze bio-economic spillovers, so I focus exclusively on Bay licenses to avoid concerns of price endogeneity.

Roughly 88% of bay shrimp landings, and 84% of total trips, happen within one of two designated seasons every year: a “Spring” open season from May 15th to July 15th, and a “Fall” open season from August 15th to November 30th (see Figure 2). To avoid anomalous harvest activity on off-season days, and as an attempt to isolate commercial activity rather than recreational, the analysis involved in this paper focuses on in-season activity only.

2.2 The Texas Shrimp License Buyback Program

Texas inshore shrimping expanded over the latter half of the twentieth century to the point of severe over-capitalization. TPWD characterized the bay shrimp fishery as simultaneously over-fished, fully exploited, and supporting more vessels than the resource could sustainably bear [McKinney, 2004]. The agency’s entry-side response came in 1995, when the Texas Legislature established a limited-entry program for the commercial bay and bait shrimp fisheries, placing a moratorium on the issuance of new licenses and fixing the pool of active shrimp fishers.

The moratorium halted further growth in capacity but could not draw down the fleet already in place, and a secondary market of license sales between fishers prevented the desired decrease in capacity. This motivated the inclusion of additional levers to retire

²Inverse demand, instrumented by weather, shows a negligibly small though statistically significant market demand elasticity of -11 , implying a 10% shock to supply leads to a 0.9% price change. The residual demand facing an individual vessel is more elastic still: see Table A.3




Figure 2: Daily Landings Averaged Across 2008 to 2023 (7-day rolling mean). Spring and Fall open-seasons illustrated above, from May 15th to July 15th and August 15th to November 30th, respectively.

licenses from the now-frozen pool. Policymakers noticed, however, that the overcapitalized fishery is also quite economically fragile. National Marine Fisheries Service economic data show that although U.S. shrimp consumption has more than quadrupled between 1984 and 2023, that demand was met almost entirely by cheaper farmed imports: the Gulf fishery’s share of the domestic market fell to roughly 4.5% by 2023, and the inflation-adjusted ex-vessel price of Gulf shrimp declined from over \$6 per pound in 1984 to under \$2, its lowest level on record [Liese et al., 2026]. Additional financial pressures include volatile diesel costs, an aging fleet, and recurrent threats from tropical storms, leaving the industry economically fragile. For this reason, TPWD deliberately considered voluntary instruments alternatives that would force vessels out of production or restrict their economic capabilities. As Coastal Fisheries Division director Larry McKinney explained, the buyback was chosen as the “... least-disruptive management tool available,” precisely because the goal was to stabilize the fishery “... socially, economically and biologically,” without further damaging an already precarious market [McKinney, 2004].

The buyback program was implemented in 1997 as a reverse auction in which bay and bait shrimp license-holders bid an asking price to voluntarily forfeit a license. At least one round (at most two) of the buyback has been run annually for every year except for 2023, with the number of purchased licenses limited by a fixed-round budget. Auctions commenced after the implementation of the 1995 moratorium on new license issuance, ensuring a fixed pool of potential entrants into perpetuity or until the policy is changed. All licensed shrimpers are eligible to participate regardless of recent harvesting activity. Of the 504 distinct license-holders who submitted bids between 2008 and 2023, 218 recorded no shrimp landings anywhere in the ticket data over that entire sixteen-year period. Inactivity is even more pronounced at the bid level: because anyone may bid repeatedly, approximately 66% of all asks during this period sought payment for licenses associated with no contemporaneous fishing activity, and therefore requested non-additional retirements.

In the 26 years since its inception, the program has run over 41 rounds, retiring 2,120




Figure 3: Texas Bay-Shrimp Landings (million lbs tails).

of a pre-moratorium pool of 3,231 bay and bait shrimp licenses (65.56% of total licenses at the onset), at a total cost exceeding \$15 million. Effort reduction has been successful: Texas bays recorded roughly 10,000 trawl-days recorded, about half the 20,000 recorded two decades prior [Texas Parks and Wildlife Department, 2007]. Recorded landings since 2005 have fluctuated between 0.5 and 3.5 million lb tails, compared to between roughly 9 and 15 million before the buyback (see Figure 3). 92% of vessels held both Bait and Bay licenses at the beginning of the program. Vessels are allowed to hold multiple of the same license (e.g. two bay licenses), in theory to increase the limit to which they can fish, though only 0.08% of in-season trips went over the per-license daily limit of 600 from 2008 to 2023, suggesting vessels carry this license as a transferable asset in the secondary market rather than to augment operational fishing capacity.

The auction operates as follows. In late January, license-holders choose whether or not to participate in the auction. In the auction season (typically end of peak season in mid July or early November), entered license-holders can choose to bid an asking price (\$/license) at which they would be willing to permanently forfeit their fishing rights. Participants are not informed of who they are bidding against. TPWD receives these bids and calculates a license-specific estimated market value (EMV), based on the age of the license and the length of the associated vessel, which serves as a reserve price. Bidders are aware that their EMV is conditioned on these variables, but don't know the specific formula.³ Bids are sorted by

³The formula used to calculate EMV was adjusted at three points, at rounds 3, 6, and 34. Each new formula increased the EMV for every license permissible in the auction to elicit more participation. This appears to have had a weak effect on increasing bid amounts: see Panel (A) of Figure 4.

their gap (EMV - Bid), and TPWD accepts bids beginning with the largest gaps first until either their budget for that round, or the set of positive gaps, is exhausted.⁴

License-holders whose bids are accepted may subsequently rescind their offers. This occurred for 376 of 2,619 accepted bids, or about 14%. Rescission carries no penalty for future participation, and returning bidders are treated no differently from first-time entrants. This option may reflect changes in willingness to sell between auction entry and final acceptance, but it may also facilitate strategic learning about the acceptance rule. Consistent with the latter interpretation, 38 of the 42 vessels that had a bid accepted, rescinded, and then returned in the following round and submitted a bid at least as high as before. By contrast, among vessels whose bids were rejected and who returned in a later round, a majority subsequently lowered their asks (see Panel (B) of Figure 4).

TPWD has run at least one round of the auction every year since 1997, except in 2023. Rounds were typically held at end of season, though many years held an additional round at the end of the spring season.⁵ The bidder pools started large, fluctuating between roughly 75 and 150 bidders, until an eventual decline after round 18 (see Panel (C) of Figure 4). The types of entrants has also changed, with the average vessel length associated with a submitted bid increasing from 35 feet in round 1 to 45 feet by round 40 (see Panel (D) of Figure 4). This compositional change may reflect a thinning of the marginal vessels in operation, capitalization on a thinner rival vessel pool (also known as “capital stuffing”), or general industry trends.

3 Data

TPWD provided vessel-level trip-ticket records from 2008 to 2023, reporting each commercial landing by vessel identifier, date, species (bay or bait shrimp), bay system, landed weight of caught shrimp, and shrimp prices by bay system. I refer to these records throughout as the ticket data. While 8 bay systems exist in the dataset, the majority of shrimp landings occur in 5 bay systems (92%). I focus exclusively on the Galveston, Matagorda, Aransas, Corpus Christi, and San Antonio Bay systems (see Figure 1). The trip tickets serve two roles, as both a bay-day harvest series for the stock and production models and vessel-level daily location choices, while the auction records supply round-by-round participation and bid outcomes. Unfortunately, trip-ticket data at the vessel level only exists in years 2008 to 2023 (rounds 20 to 40), which excludes some of the years with the largest bidder pools.

I use the TPWD boat/motor ownership inquiry application (TORA) to merge additional vessel-specific characteristics for vessels registered at the state level [[Texas Parks and Wildlife Department, 2026](#)]. Vessels weighing more than 5 net tons only require U.S. coast guard registration and do not appear in the TORA inquiry system [[United States Coast Guard, 2026](#)]. Many vessels have both, and are merged using an ID crosswalk filling roughly 79% of all vessels appearing in ticket and bid data. Unfortunately, the only vessel-specific characteristics available in both systems are year built and vessel length, though I do include addresses associated with vessel registration from TORA to approximate travel costs from

⁴Rejected license-holders are not informed whether their bid was declined due to an insufficient gap relative to rivals or because the round’s budget was exhausted before their bid was reached.

⁵Years with two rounds held included 1997-99, 2002-05, 2007-11, and 2015.




Figure 4: Summary Figures on the Buyback Auction

Note: Four figures above summarize auction activity. Panel (A) shows placed bids per round, with vertical lines denoting when TPWD changed the estimated market value of a license used as the reserve price, i.e. the maximum accepted price of a license based on vessel length and license age. Panel (B) shows difference in bid amount conditional on prior-round bid. Panel (C) shows bidder pools over time. Panel (D) shows boat length over time, increasing slowly.

home-to-bay system centroids.

I combine vessel characteristics with a historical record of auction participation, provided by TPWD, recording thousands of bids over 41 rounds from 1997 to 2024. I restrict the sample to rounds that can be merged with active landed tickets, reducing this sample to rounds 20 through 40 (years 2008 to 2023). As mentioned in the previous section, auction outcomes are dictated by TPWD’s set auction rules. Using a formula to calculate estimated market value (EMV) of the corresponding license, they ‘grant’ bids in descending order of EMV minus bid until their budget for that round would be exhausted. Bidders are allowed the opportunity to ‘accept’ this transaction or to ‘rescind’ their option to sell their license. Finally, I merge into this dataset time-varying covariates that influence vessel productivity to discipline empirical models. Average daily wind speed, measured in meters per second, and wave height, measured in meters, are taken from the closest buoy to the centroid of the five bay systems [NOAA National Data Buoy Center, 2026]. Gulf-specific diesel prices per gallon are taken from the FRED database [Federal Reserve Bank of St. Louis, 2026]. Summary statistics are reported in Table 2, organized in panels by the unit of observation.

4 A Dynamic Model of Common-Pool Exploitation

This section develops the structural framework used to recover the continuation value of holding a shrimp license. The empirical strategy proceeds in three steps. First, I estimate harvest production and stock depletion using vessel-level trip-ticket data, isolating biological depletion from congestion. Second, I estimate vessels’ daily participation and bay-choice probabilities and use these, together with estimated state transitions, to recover structural payoff parameters and forward-simulate continuation values. Third, I merge those continuation values with the auction records to estimate how the value of remaining in the fishery affects the decision to bid. These estimated objects are then used in Section 6 to simulate the buyback’s participation and effort-offset effects over the program’s full horizon.

4.1 Model Overview

I model the Texas inshore shrimp fishery as an infinitely repeating finite-horizon dynamic game in which heterogeneous vessels make daily participation and location decisions across multiple bay systems as a best response to rival behaviors. The framework extends Huang and Smith [2014], who use methods proposed by Aguirregabiria and Mira [2002, 2007] to estimate efficiency costs of common-pool exploitation in the North Carolina shrimp fishery using a single estuary with a binary fish/don’t-fish decision, to a spatial ‘patchy’ setting with five bay systems (Galveston, Matagorda, San Antonio, Aransas, and Corpus Christi) and endogenous location choice, following related work on common-pool resource exploitation across a patch environment [Hotz and Miller, 1993, Smith, 2005, Sanchirico and Wilen, 2020].

Texas inshore fisheries are open from May 15th to November 30th, with an inter-season regulatory closure from July 15th to August 15th. Because shrimp are highly fecund and migrate offshore to spawn, the cross-year recruitment overfishing problem can be ignored, and I only need to consider overfishing’ impact on biomass growth within a given year. The




Figure 5: Overview of the structural pipeline. Raw data objects are shaded gray, estimated objects blue, simulated objects green, and results amber.

Table 2: Summary statistics

Variable	Mean	SD	Min	Max	N
Panel A. Vessels (pooled over all auction rounds)					
Vessel length (ft)	37.57	11.10	15.00	70.00	1,493
Year Built	1980	10	1900	2020	1,571
Ever bid	0.47	0.50	0.00	1.00	1,593
Ever removed	0.35	0.48	0.00	1.00	1,593
<i>Active in tickets (2008–2023)</i>					
Total landings (1,000 lbs)	46.38	67.76	0	487.62	557
Bays fished	1.57	0.88	1.00	5.00	557
Years active	6.46	5.32	1.00	16.00	557
Panel B. Bay-day covariates (2008–2023)					
Vessels on water (per day)	10.51	13.13	1.00	90.00	9,694
Daily harvest (lbs)	2,166.50	4,261.29	0.19	90,024.25	9,694
Cumulative harvest (1,000 lbs)	86.08	137.96	0.00	976.00	9,694
Ex-vessel price (\$/lb)	3.19	2.45	0.03	61.23	9,694
<i>Time-varying only (Averaged daily, statewide / single buoy)</i>					
Wind speed (m/s)	5.36	1.83	1.21	15.12	9,694
Wave height (m)	0.76	0.34	0.14	4.51	9,694
Diesel (\$/gal)	2.37	0.79	0.85	4.57	9,694
Panel C. Bay-license bids (all rounds)					
Bid (asking price, 1,000\$)	10.78	12.56	0.10	200.00	2,515
EMV reserve (1,000\$)	7.23	1.52	1.88	10.07	2,515
Gap (EMV – bid, \$)	-3.55	12.03	-196.71	8.21	2,515
Granted	0.53	0.50	0.00	1.00	2,481
Removed	0.44	0.50	0.00	1.00	2,496
Rescinded	0.08	0.27	0.00	1.00	2,468

Note: Panel A covers all vessels appearing in bay tickets or bay bids. Length and year built come from TORA/USCG registration records and are missing where a vessel could not be matched. The italicized rows describe only vessels with recorded bay landings. Panel B is bay-system-by-day, in-season, 2008–2023, restricted to days with at least one active vessel. Panel C covers bay-license bids across all 41 rounds (1997–2024); EMV is computed from the TPWD 55–1000 matrix (\$55/ft plus a license-history value) with era multipliers (+16% in rounds 4–5; +16% and +10% from round 6) and a 7.5% premium when two licenses were tendered, crediting 5 years of history from round 6 on. 45 non-positive bid(s) and corresponding EMVs and gaps are dropped as evident misunderstandings of the reverse-auction rules or faulty entries.

annual cycle simplifies the stock externality problem: within-season depletion by the fleet reduces current-year stocks, but through exogenous recruitment, the stock renews partially

during each one-month inter-season closure, and fully renews during each six-month Winter closure.

Each day $t = \{1, \dots, T = 170\}$ within season τ , vessel i observes a common state vector $\mathbf{s}_t$ comprising bay-specific stock indices, shrimp prices, weather conditions, diesel prices, congestion levels, and loyalty profiles across bay systems. The vessel then chooses among $J + 1$ alternatives: fish one of the J bay systems, or refrain from fishing. Vessels are forward-looking with respect to the evolution of stock, prices, and rival behavior; because any single vessel’s harvest is a negligible share of a bay’s aggregate stock, however, each treats the stock path as essentially exogenous to its own daily choice. I assume the observed fleet behavior constitutes a Markov Perfect Equilibrium (MPE).

The model is estimated in two stages. In the first stage, I recover the production technology, latent stock dynamics, conditional choice probabilities (CCPs), and state variable transition processes. In the second stage, I use the first-stage CCPs and transition estimates to simulate continuation values forward and recover the profit parameters via pseudo-maximum likelihood. The CCP approach avoids solving the full dynamic programming problem, which would be computationally intractable given the continuous state space and hundreds of strategic agents.

Two features of this environment are worth clarifying. First, the model delivers two distinct constructs with different objectives. The first includes structural payoff parameters $\boldsymbol{\theta}$ that describe preferences, such as how a vessel’s owner trades off expected revenue, stock, loyalty, travel cost, and weather in daily location choices, recovered by matching observed daily choices. The second is continuation value CV_i , the expected discounted value of holding a license and continuing to fish, which is a separate quantity generated by forward simulating the daily choice problem dictated by $\boldsymbol{\theta}$ that license-holders weigh against a buyback offer. The second feature is to clarify that each vessel is atomistic. Because own-harvest is negligible relative to a bay’s stock, an individual’s daily choice is effectively myopic, even as the fishery remains fully dynamic in the aggregate and combined harvest depletes the common stock. The dynamic common-pool externality the buyback generates therefore does not live in any single vessel’s daily problem, but rather in the aggregate stock and, through it, in the continuation value of the license (see [A.1](#)).

4.2 Production Function and Stock Dynamics

4.2.1 Harvest Production

I assume harvest for vessel i in bay j on day t follows a restricted Cobb-Douglas production function:

$$h_{ijt} = q_i \cdot \exp(\gamma(N_{jt} - 1) + \delta'W_t) X_{jt} \exp(\varepsilon_{ijt}^h) \quad (1)$$

where q_i is vessel-specific catchability capturing unobserved fishing technology, captain skill, and experience; N_{jt} is the count of vessels fishing bay j on day t (entered as leave-one-out to avoid own-vessel contributions to congestion); W_t collects daily weather controls including wind speed $WSPD_t$ and wave height $WVHT_t$ measured at NOAA buoys [[NOAA National Data Buoy Center, 2026](#)], and summer regulatory closure indicators; X_{jt} is the latent stock index in bay j on day t ; and ε_{ijt}^h is an i.i.d. harvest shock with mean zero.

The parameter γ captures congestion or agglomeration effects. If $\gamma < 0$, additional vessels reduce individual harvest, consistent with negative congestion externalities as shrimp trawling disperses aggregations of stock. If $\gamma > 0$, additional vessels increase productivity, perhaps through information sharing that reduces search costs.

After log transformation, the production function becomes a linear regression:

$$\ln(h_{ijt}) = \ln(q_i) + \gamma(N_{jt} - 1) + \delta'W_{jt} + \ln(X_{jt}) + \varepsilon_{ijt}^h \quad (2)$$

where $\ln(q_i)$ is estimated with vessel effects.

If vessels observe bay-specific productivity shocks invisible to the researcher, a positive shock attracts more vessels, generating spurious positive correlation between vessel counts and harvest. OLS estimates of γ are therefore biased upward. [Huang and Smith \[2014\]](#) instrument total vessel count with a weekend indicator, shifting participation decisions through non-economic channels orthogonal to stock shocks. This is not applicable to my setting, however, as weekend (or other fishery-wide regulatory) shocks do not vary over j . Instead, I instrument bay-specific vessel counts using bay-specific Bartik/shift-share instruments.

4.2.2 Latent Stock Index

The stock is not directly observed by the researcher, but is instead treated as a latent state variable observable by fishermen and endogenously influenced by their activities. I decompose the bay-specific stock index as

$$X_{jt} = A_{jt} \cdot X_{0,\tau} \cdot G_j(t) \cdot \exp(-\phi H_{jt}^{cum}) \cdot \exp(\zeta_{jt}) \quad (3)$$

Where A_{jt} represents bay system fixed effects; $X_{0,gr}$ is the initial stock size at the beginning of year gr (captured by year fixed effects), reflecting exogenous recruitment; $G_j(t)$ is the deterministic within-season biological growth function, which is bay-specific and fixed across years; ϕ is the fishing mortality coefficient, measuring how cumulative harvest $H_{jt}^{cum} = \beta_{0j} + \sum_{s=0}^t h_j(s)$ depletes the stock relative to its biological baseline⁶; and β_{0j} is any accumulated harvest from January 1st through May 14th of that year; and ζ_{jt} is a stochastic shock to the stock index.

4.2.3 Within-Season Growth Function

The deterministic growth function $G_j(t)$ is modeled as bay-specific quartic polynomial in log-time:

$$\ln G_j(t) = \beta_{j1} \ln(t) + \beta_{j2} [\ln(t)]^2 + \beta_{j3} [\ln(t)]^3 + \beta_{j4} [\ln(t)]^4 \quad (4)$$

where t is days since May 15th (the approximate start of the harvest season). Note the first three coefficients vary by bay system while the fourth is pooled across bay systems due to power issues.

After estimating the production function, the composite residual $\varepsilon_{ijt}^h + \zeta_{jt}$ is recovered. Since ε_{ijt}^h is i.i.d. across vessels and ζ_{jt} is common to all vessels in bay j on day t , regressing the residual on bay-by-day fixed effects identifies ζ_{jt} .

⁶ H_{cum} is reset to zero at the mid-summer closure, following [Huang and Smith \[2014\]](#).

4.3 Bay System Choice

The first-stage CCP estimation recovers the probability that vessel i chooses among $J + 1$ alternatives on day t . To accommodate the participation and location decisions jointly while allowing for correlation among the fishing alternatives, I estimated a nested logit with two nests: a fishing nest containing five bay systems and a singleton “home nest”. The deterministic utility for fishing bay j is:

$$V_{ijt} = \alpha_{0j} + \sum_{k=1}^3 [\beta_{p^k} p_{jt}^k + \beta_{X^k} X_{jt}^k] + p_{jt}^{ma} + \beta_L L_{ijt} + \beta_s C_t + \beta_W W_t + \ln q_i + \tau_t \quad (5)$$

where α_{0j} is a bay-specific alternative-specific constant; p_{jt} is the shrimp price at bay j ; p_{jt}^{ma} is a 30-day weighted moving average of bay system prices; X_{jt} is the latent stock index; C_t indicates the inter-season regulatory closure; W_t includes wind speed and wave height recorded at the nearest buoy; k indexes the order of polynomials; and τ_t represent year fixed-effects. The outside option is normalized at $V_{i0t} = 0$.

Fishermen exhibit strong persistence in location choice, returning to familiar grounds even when conditions elsewhere may be more favorable [Smith, 2005, Sanchirico and Wilen, 2020]. I model this through L_{ijt} , defined as an exponentially weighted average of vessel i ’s recent history of trips to bay j :

$$L_{ijt} = a_L L_{ij,t-1} + (1 - a_L) \mathbf{1}[a_{it} = j] \quad (6)$$

where a_L is a smoothing parameter, calibrated to fit bay system choices recorded in the ticket data. This specification captures the notion that vessels develop bay-specific knowledge (bottom conditions, shrimp aggregation patterns, local market relationships, etc.) that create inertia in location choice. Under the nested logit assumption with inclusive value parameter $\lambda \in (0, 1]$, the probability that vessel i chooses bay j is:

$$\Pr(j|i, t) = \frac{\exp(V_{ijt}/\lambda)}{\sum_{k \in \mathcal{J}_F} \exp(V_{ikt}/\lambda)} \cdot \frac{\exp(\lambda I_{it})}{1 + \exp(\lambda I_{it})} \quad (7)$$

Where $\mathcal{J}_F$ is the fishing nest and $I_{it} = \ln \sum_{k \in \mathcal{J}_F} \exp(V_{ikt}/\lambda)$ is the inclusive value of the fishing nest. The probability of staying home is $\Pr(0|i, t) = 1/(1 + \exp(\lambda I_{it}))$.

4.4 State Variable Transitions

With the continuation values expressed as functions of the first-stage CCPs, estimating the structural parameters requires evaluating what those continuations would be at states the vessel has not yet visited. Following Bajari et al. [2007], I simulate the exogenous state vector forward from each observed date t to the end of the season using estimated transition processes for each state component. Shrimp prices follow a bay-specific VAR(4):

$$\ln \mathbf{p}_{t+1} = \boldsymbol{\mu}_p + \sum_{\ell=1}^4 \boldsymbol{\Phi}_\ell \ln \mathbf{p}_{t-\ell+1} + \boldsymbol{\varepsilon}_t^p, \quad \boldsymbol{\varepsilon}_t^p \sim N(\mathbf{0}, \boldsymbol{\Sigma}_p), \quad (8)$$

where $\mathbf{p}_t$ is the J -vector of bay prices and $\varepsilon \sim N(0, \Sigma_p)$. Weather variables (wind speed and wave height) follow a VAR(3):

$$\ln \mathbf{W}_{t+1} = \boldsymbol{\mu}_W + \mathbf{A}_1 \ln \mathbf{W}_t + \mathbf{A}_2 \ln \mathbf{W}_{t-1} + \mathbf{A}_3 \ln \mathbf{W}_{t-2} + \boldsymbol{\varepsilon}_t^W, \quad \boldsymbol{\varepsilon}_t^W \sim N(\mathbf{0}, \boldsymbol{\Sigma}_W), \quad (9)$$

where $\varepsilon_t^w \sim N(0, \Sigma_w)$ and the covariance allows correlation between wind speed and wave height within a buoy. Diesel prices follow a random walk in logs,

$$\delta_{t+1} = \delta_t + \varepsilon_t^\delta, \varepsilon_t^\delta \sim N(0, \sigma_\delta^2) \quad (10)$$

and stock shocks ζ_t follow a bay-specific AR(2):

$$\zeta_{j,t+1} = \rho_{1j} \zeta_{jt} + \rho_{2j} \zeta_{j,t-1} + \varepsilon_{jt}^\zeta, \quad \boldsymbol{\varepsilon}_t^\zeta \sim N(\mathbf{0}, \boldsymbol{\Sigma}_\zeta), \quad (11)$$

These transition processes are estimated using the full panel.

4.5 Equilibrium and Continuation Values

4.5.1 Markov Perfect Equilibrium

Each day t , vessel i chooses action $a_{it} \in \{0, 1, \dots, J\}$ to maximize expected discounted utility over the remaining season:

$$\max_{a_{it}} \{u_i(a_{it}, \mathbf{s}_t; \boldsymbol{\theta}) + \xi_{ia_{it}t} + \beta_d \mathbb{E}_t[V_i(\mathbf{s}_{t+1})]\} \quad (12)$$

where $u_i(\cdot)$ is the per-period structural payoff, s_t denotes exogenous state variables, $\xi_{ia_{it}t}$ is a private preference shock within the nest, distributed Type I Extreme Value, and β_d is the daily discount factor. The per-period structural utility for choosing bay j by vessel i is:

$$\begin{aligned} u_i(j, \mathbf{s}_t; \boldsymbol{\theta}) = & \kappa_{0j} + \alpha(p_{jt} \hat{h}_{ijt} - \bar{r}_i) \\ & + \beta_X X_{jt} + \beta_L L_{ijt} + \beta_d d_{ij} \delta_t + \beta'_W \mathbf{W}_t + \beta_\delta \delta_t + \beta_w \text{wknd}_t + \beta_C C_t + \beta_\ell \ell_i + \beta_{\ell\ell} \ell_i^2 \end{aligned} \quad (13)$$

Here, κ_{0j} represents the five alternative-specific constants of the bay system, and α represents the marginal utility of expected revenue, in the demeaned revenue of the vessel $(p_{jt} \hat{h}_{ijt} - \bar{r}_i)$ ⁷. The outside option is normalized to zero ($u_i(0, s_t; \boldsymbol{\theta}) = 0$), and the structural parameter vector is $\boldsymbol{\theta} = (\alpha, \beta_X, \beta_L, \beta_d, \beta_W, \beta_\delta, \beta_w, \beta_C, \beta_\ell, \beta_{\ell\ell}, \{\alpha_{0j}\}_{j=1}^5, \lambda)$. d_{ij} represents a proxy for distance, measured between a bay system's centroid and the most recently reported address associated with a license on TORA. Distance only enters as an interaction with diesel prices, δ_t , which also enters as its own marginal effect on utility. ℓ represents

⁷Revenue enters this way because catchability enters revenue multiplicatively, and the revenue level is collinear with q_i ; the demeaning within the vessel identifies α from variation around each vessel's own mean rather than the confounded cross-vessel level.

vessel length⁸, $wknd_t$ is an indicator variable set to 1 for Saturdays and Sundays, and λ is a nested inclusive value. A Markov Perfect Equilibrium is a profile of strategy functions $\{\varphi_{it}^*(\mathbf{s}_t, \xi_{it})\}$ such that each vessel's strategy is a best response to the strategies of all other vessels:

$$\varphi_{it}^*(\mathbf{s}_t, \xi_{it}) = \arg \max_{a_{it}} [V_{it}^\varphi(a_{it}, \mathbf{s}_t, \xi_{it})] \quad (14)$$

where the value function accounts for the equilibrium play of all other vessels both in the current period and in all future periods. This equilibrium concept imposes that decision rules and beliefs about others' behavior are mutually consistent.

4.5.2 Forward Simulation

Under the Type I Extreme Value assumption, CCPs contain sufficient information to recover continuation values non-parametrically, avoiding backward induction entirely. Specifically, the expected value of choosing optimally prior to the realization of private shocks has the closed form:

$$V_i(\mathbf{s}_t) = \ln \sum_{j=0}^J \exp(u_i(j, \mathbf{s}_t; \theta) + \beta_d \mathbb{E}[V_i(\mathbf{s}_{t+1}) | a_{it} = j]) + \gamma_E \quad (15)$$

where γ_E is Euler's constant. Because the outside option is normalized, the log-probability of staying home provides a nonparametric estimate of the continuation value:

$$-\ln(\hat{\sigma}_{i0t}) = \ln(1 + \exp(\lambda I_{it})) \quad (16)$$

The nested logit introduces the inclusive value parameter λ , so the continuation value of the fishing nest is:

$$CV_{it} = \lambda \ln \sum_{k \in \mathcal{J}_F} \exp(\hat{V}_{ikt}/\lambda) \quad (17)$$

To evaluate continuation values at candidate parameter values for states the vessel has not visited, I simulate the state vector forward. On each simulated day, the procedure:

1. Draws next-period state variables from the estimated transition processes (prices, weather, diesel, stock shocks).
2. Updates the stock index incorporating harvest feedback: the deterministic stock evolves via $G_j(t)$, while cumulative harvest depletes the stock through $\exp(-\phi H_{jt}^{cum})$.
3. Computes nested logit choice probabilities for all vessels given the current state, yielding predicted vessel counts by bay.

⁸Length, as well as distance, are only available as the most recent recording available on the TORA and/or USCG database as of January 2025. If a license-holder changed addresses, or if the associated vessel has been capitalized since their appearance in the data, it is not observable. Bid data includes a contemporaneous record of vessel length, which I default to when available. These should be seen as noisy proxies for distance and length, respectively.

4. Updates loyalty based on realized (simulated) choices.
5. Iterates to a fixed point for vessel counts within each day to ensure consistency between predicted congestion and the choice probabilities that depend on it.
6. Accumulates the discounted log-sum contribution from each simulated day.

The simulation runs from each observed date t forward to the end of the season, and the discounted sum of daily log-sum values gives the continuation value. The daily discount factor β_d is set to 0.999.

4.5.3 Continuation Value Computation

The structural parameters are estimated via pseudo-maximum likelihood:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i,t} \sum_{j=0}^J \mathbf{1}[a_{it} = j] \cdot \ln \Pr(j|\mathbf{s}_t; \theta) \quad (18)$$

where the choice probabilities are constructed from the choice-specific values $\tilde{V}_{ijt}(\theta) = u_i(j, \mathbf{s}_t; \theta) + \beta \cdot \bar{C}V_{it}$, with $\bar{C}V_{it}$ denoting the simulation-averaged continuation value. The advantage to this approach is that θ only enters through the per-period utility $u_i(\cdot)$, while continuation value is estimated nonparametrically from the CCPs and does not depend on θ , following from the atomicity of the individual vessel⁹. Because a single vessel's harvest is negligible relative to a bay's stock, the continuation value on the right-hand side of (15) is invariant to that vessel's own current choice: fishing bay j today rather than staying home leaves the aggregate stock path, and option value of continued operation, essentially unchanged. The continuation term is thus common across a vessel's alternatives on a given day, so it drops out of the conditional choice probabilities and does not confound the flow parameters θ . This is the finite-dependence logic proposed by Arcidiacono and Miller [2011], specialized to an atomistic common-pool setting. The one channel through which a vessel's own choice does move its future state is loyalty — fishing bay j today raises its propensity to return — but a single day's update is second-order over a season of $T = 170$ days, and I treat the continuation as invariant to it.

For the buyback analysis in subsequent sections, I convert the season-level continuation value from logsum units to dollars using the structural revenue coefficient from Equation (13): $CV_i^{\$} = CV_i/(\alpha \times 1000)$. The vessel-specific continuation value accounts for heterogeneity in catchability q_i , home port location, and loyalty profiles. Under an infinite-horizon interpretation with a 7.5 to 9% annual discount rate, the present value of the license is approximately $CV_i^{\$}/(1 - \delta)$, where $\delta = 0.93$ is the annual discount rate¹⁰.

⁹This assumption appears in the data: the median vessel accounts for under one-percent (0.88%) of its bay-year harvest and the effective number of vessel per bay-year — the inverse Herfindahl index — is roughly thirteen. This holds for all bay systems, even though this paper only focuses on the five most active bay systems, where atomicity is held more strongly. No single vessel materially moves the stock it responds to (see Appendix A.1)

¹⁰This annual discount rate is embedded at a different horizon compared to the daily discount factor set earlier 0.999 daily discount rate used for the forward-simulation in Section 4.5.2.

5 Estimation Results: Fishery Dynamics

5.1 Production Function Estimates

Table 3 reports estimates of the production function in Equation (2). Because including vessel i 's own presence in the congestion measure would mechanically generate a positive correlation between vessel counts and harvest, I use leave-one-out vessel counts ($N_{jt} - 1$) throughout. Column (1) shows the results of regressing log harvest on congestion, weather variables, cumulative season harvest, regulatory closure (July 15th to August 15th), Bay $\times$ time polynomials, and year, vessel, and bay fixed effects.

The coefficient on vessel count is positive and significant (0.017), suggesting agglomeration effects dominate congestion. However, as discussed in Section 4, this estimate is biased if unobservable stock shocks attract vessels to productive bays. While Huang and Smith [2014] use weekend closures as an exogenous instrument, this does not work in my setting as closures do not vary across bay systems and would not instrument bay-specific vessel counts in a model with bay fixed effects. I construct bay-specific Bartik/shift-share instruments from historic activity. For each vessel m and bay j , I record the historical fraction of m 's trips that go to bay j (computed while leaving own year out), with vessels that never fished a bay getting $s_{mj} = 0$. I then generate shifts, $\hat{P}_{mt}$ using a vessel-level participation logit for whether vessel m is active on day t using only common, exogenous drivers — wind speed, wave height, diesel prices, a time polynomial, and weekend, closure, and year dummies. These shifts are each vessel's predicted probability of being out on the water that day for reasons unrelated to any particular bay's conditions. Finally, I combine these to create the Bartik instruments: the share-weighted sum of rivals' predicted activity:

$$Z_{ijt} = \sum_{m \neq i} s_{mj} \hat{P}_{mt} \quad (19)$$

Columns (2) and (3) detail the results of the 2nd and 1st stage estimation of the model using Bartik instruments for bay-specific vessel counts. After instrumenting, the coefficient associated with congestion is not statistically different from zero. This is not surprising: Texas Gulf Coast bay systems are large (summing to a total of 4,363 km²) and contain multiple industries within, including crab and oyster fisheries.

5.2 Stock Dynamics

The latent stock index is recovered as the systematic component of the production residual: from equation (3), $\ln X_{jt}$ decomposes into a bay baseline A_j , an annual level $X_{0,\tau}$ (the year effect, capturing exogenous recruitment), a bay-specific within-season growth path $G_j(t)$, the cumulative-harvest depletion ϕH_{jt}^{cum} , and a serially correlated shock ζ_{jt} . Three features merit description. First, the deterministic path rises early in the season and bends downward as cumulative removals accumulate; the gap between the with- and without-harvest paths in Figure 6 is the depletion harvest creates. Second, that depletion is economically large — a one-million-pound increase in cumulative seasonal harvest lowers the stock index by roughly 47% ($\hat{\phi} = -0.64$; Table 3, column 2) — and it is the biological force the buyback seeks to relax. Third, the stock shocks ζ_{jt} are mean-zero and modest but moderately persistent

Table 3: Production Function Estimates

	(1) OLS $\ln h$	(2) IV (2nd Stage) $\ln h$	(3) IV (1st Stage) $n_vessels$
<i>Congestion</i>			
Count of Vessels (minus one)	0.0172*** (0.0002)	−0.0004 (0.0008)	
Predicted rivals (shift-share IV)			0.8811*** (0.0090)
<i>Environmental controls</i>			
WSPD	0.0047* (0.0026)	0.0089** (0.0032)	0.5604*** (0.0401)
WVHT	−0.0397** (0.0129)	−0.1225*** (0.0182)	−0.3915 (0.2287)
<i>Economic controls</i>			
Cumulative harvest (millions of lbs)	−0.3376*** (0.0301)	−0.6429*** (0.0298)	−19.7550*** (0.3621)
Summer closure	−0.4619*** (0.0127)	−0.7803*** (0.0188)	−7.0375*** (0.1884)
Bay $\times$ time polynomial	✓	✓	✓
Year FE	✓	✓	✓
Vessel FE	✓	✓	✓
Bay FE	✓	✓	✓
Observations	101,818	101,818	101,818
Adj. R^2	0.359	0.320	0.652
Within R^2	0.122	0.076	0.404
RMSE	0.907	0.931	11.700
<i>IV diagnostics</i>			
First-stage F		9,560.67	
Wu-Hausman p -value		0.000	

Notes: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Standard errors in parentheses. The endogenous regressor is the leave-one-out count of rival vessels, instrumented with a shift-share prediction of rivals.




Figure 6: Logged Stock Indices by Bay System

Note: Deterministic stock indices in blue; deterministic stock overlaid in red. Stock estimated only within seasonal harvest days, May 15th to November 30th, with inter-season closure July 15th to August 15th

(per-bay AR(2); Table 4), so day-to-day deviations from the deterministic path carry over for several days rather than immediately reverting to zero. The growth-polynomial, bay, and year-effect estimates underlying the index are reported in Appendix A.2.

5.3 State Transition Fit

Following Bajari et al. [2007], I simulate the exogenous state vector $\mathbf{s}_t$ forward from each observed date to the end of the season, drawing each component from an estimated transition process. The components are bay-level log shrimp prices p_{jt} , wind speed $WSPD_t$ and wave height $WVHT_t$, the diesel price d_t , and the stock index X_{jt} . Dickey-Fuller tests reject the unit-root null ($p \approx 0$) for every process except diesel; I model log prices as a VAR(4) process, weather as a VAR(3) process, bay-specific shocks as an AR(2) process, and diesel

prices, which retain a unit root, enter as a random walk with drift. Table 4 reports the estimates.

Table 4: State Variable Transition Processes

	Persistence (own, $\sum$ lags)	Cross-bay ($\sum$ off-diag)	Innovation SD
<i>Bay-level log price (VAR(4))</i>			
Aransas Bay	0.929	0.319	0.286
Corpus Christi Bay	0.897	0.444	0.364
Galveston Bay	0.848	0.275	0.223
Matagorda Bay	0.758	0.600	0.436
San Antonio Bay	0.889	0.339	0.301
<i>Weather (VAR(3))</i>			
Wind speed	0.461	0.181	0.290
Wave height	0.653	0.428	0.294
<i>Diesel (random walk with drift)</i>			
Log diesel price	1.000	—	0.0294
<i>Stock innovation ζ (AR(2), by bay)</i>			
Aransas Bay	0.828	—	0.378
Corpus Christi Bay	0.681	—	0.552
Galveston Bay	0.755	—	0.333
Matagorda Bay	0.609	—	0.411
San Antonio Bay	0.654	—	0.455

Notes: Price and weather are vector autoregressions estimated jointly across bays/variables with a constant; “Persistence” sums the own-series coefficients across all lags and “Cross-bay” the absolute off-diagonal coefficients across all lags. Innovation SD is $\sqrt{\text{diag}(\Sigma)}$. The stock innovation ζ follows a per-bay AR(2) with cross-bay innovation covariance. All mean-reverting processes are stationary: the companion-matrix spectral radius is 0.964 for the price VAR and 0.619 for the weather VAR, and the per-bay ζ AR(2) roots have maximum modulus 0.713–0.866; every value lies strictly inside the unit circle. Diesel is modeled as a random walk with drift after first-difference stationarity tests (unit root by construction).

5.4 Stage 1 Conditional Choice Probabilities

Table 5 reports the first-stage nested logit. It’s worth mentioning that these results are not meant to carry a reduced-form interpretation, but simply as a tool for simulating vessel-bay-day choices for estimating continuation values in the second stage. Due to computation constraints, a 40% sample is run from the almost one million choice occasions that arise in the data: hundreds of vessels choose whether or not to fish one of five bay systems over the course of 170 in-season days and 16 years ¹¹

¹¹Sample restricted in-season to vessels that are active to minimize chances of modeling choice occasions from fishermen who have retired their license outside of the auction. This matches the sample restriction in

The inclusive-value parameter is 0.542 and precisely estimated, lying inside the unit interval consistent with random-utility maximization and supporting the fish/home nest. Among the common parameters, stock and ex-vessel price both enter with positive first-order coefficients, so a vessel is more likely to fish a bay when its stock and price are higher, with the cubic terms allowing those marginal effects to bend over their ranges. The loyalty term is large and tightly estimated (4.34), reflecting strong persistence in where a vessel fishes from one period to the next. The alternative-specific coefficients track physical conditions: wave height lowers participation in all bays, and the mid-summer closure sharply reduces it everywhere. Catchability enters with mixed signs: productive vessels prefer to fish in the San Antonio Bay system, while less productive vessels prefer Corpus Christi and Matagorda Bay.

5.5 Structural Profit Parameters

The second stage recovers the structural payoff parameters θ by pseudo-maximum likelihood, using the first-stage conditional choice probabilities and state-transition estimates to forward-simulate each vessel’s choice-specific value function. Because θ enters the per-period payoff linearly while the continuation value is pinned down by the first-stage CCPs, the pseudo-likelihood is maximized without re-solving the dynamic program. Table 6 reports the estimates.

The parameters are economically sensible and precisely estimated. The revenue coefficient is positive ($\hat{\alpha} = 0.097$), confirming that vessels value expected profit as the model requires. Stock enters strongly and positively (1.20): a healthier bay raises the utility of fishing it. This is the structural counterpart to the stock’s role in the production function and the mechanism through which the buyback’s biological spillover ultimately operates. Loyalty is large (2.84), echoing the strong location persistence recovered in the first stage. The cost-side terms all carry the expected signs—wind speed, wave height, the weekend indicator, and the summer closure each reduce the utility of fishing, with the closure the single largest deterrent (-1.18)—and utility declines with length, with the quadratic is insignificant, so the length effect is essentially linear-negative. This last point is consistent with the idea that larger vessels have more variable costs for each trawl day [Huang and Smith, 2014]. The inclusive-value parameter is $\hat{\lambda} = 0.44$, comfortably inside the unit interval, which both validates the nested-logit structure (consistency with random-utility maximization) and indicates meaningful correlation in the unobserved determinants of bay choice within the fishing nest.

One coefficient merits comment: the diesel *level* enters positively (0.09) even though the distance $\times$ diesel interaction is negative, as expected for a travel cost. The level term most plausibly absorbs within-season seasonality in diesel prices that co-moves with the peak harvest season rather than a genuine positive cost effect; the cost channel of fuel is identified through the interaction with distance.

the entry modeling in Section 5.6.

Table 5: Stage 1: Nested-Logit Location Choice

	Galveston	Aransas	Corpus Christi	Matagorda	San Antonio
<i>Panel A. Common parameters (generic across alternatives)</i>					
Price			0.194***	(0.045)	
Price ²			−0.104***	(0.011)	
Price ³			0.008***	(0.001)	
Price MA(30)			0.082***	(0.008)	
Stock			0.839***	(0.014)	
Stock ²			−0.086***	(0.002)	
Stock ³			0.001***	(0.000)	
Loyalty			4.349***	(0.074)	
<i>Panel B. Alternative-specific parameters (relative to home)</i>					
Intercept	−6.181*** (0.134)	−6.063*** (0.271)	−4.049*** (0.166)	−4.992*** (0.172)	−6.427*** (0.329)
Wind speed	−0.038*** (0.008)	0.011 (0.023)	−0.035 (0.023)	0.028* (0.013)	−0.009 (0.018)
Wave height	−0.259*** (0.044)	−0.517*** (0.133)	−0.364** (0.129)	−0.488*** (0.074)	−0.457*** (0.099)
Weekend	−0.135*** (0.017)	−0.118* (0.052)	0.080 (0.049)	−0.188*** (0.029)	−0.257*** (0.038)
Mid-Summer closure	−0.708*** (0.022)	−2.258*** (0.122)	−0.789*** (0.068)	−1.733*** (0.056)	−2.493*** (0.106)
$\ln q_i$	0.023 (0.017)	0.049 (0.051)	−0.576*** (0.022)	−0.195*** (0.024)	0.167** (0.062)
Year × bay effects	Yes (2008 base)				
Inclusive value (λ)	0.542*** (0.011)				
Log-likelihood	−96,777				
Choice occasions	222,663				

Notes: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. First-stage nested-logit location choice; home is the reference alternative and Panel B parameters are relative to it. The five bays form a single fishing nest with inclusive-value parameter $\lambda \in (0, 1)$, consistent with random-utility maximization. Price and stock enter as third-order polynomials; year-by-bay effects are estimated but suppressed. Sample: daily choice occasions, 2008–2023.

Table 6: Structural Profit Parameters (Second-Stage Estimates)

	Coefficient	Std. error
<i>Panel A. Utility parameters</i>		
Revenue (α)	0.097***	0.007
Stock	1.199***	0.004
Loyalty	2.839***	0.025
Distance $\times$ diesel	-0.007*	0.003
Wind speed	-0.040***	0.007
Wave height	-0.248***	0.007
Diesel	0.090***	0.004
Weekend	-0.179***	0.004
Summer closure	-1.182***	0.005
Vessel length	-0.165***	0.020
Vessel length ²	0.034	0.020
<i>Panel B. Nesting</i>		
Inclusive value (λ)	0.440***	0.011
<i>Panel C. Alternative-specific constants (rel. home)</i>		
Galveston	-4.670***	0.086
Aransas	-5.365***	0.088
Corpus Christi	-5.593***	0.089
Matagorda	-5.119***	0.087
San Antonio	-4.963***	0.086
Log-likelihood	-247,367	
Choice occasions	529,743	

Notes: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Second-stage nested-logit structural estimates; home is the reference alternative. The sample is daily choice occasions for vessels active in a given year (at least one landing that year), 2008–2023.

Standard errors are second-stage Hessian errors, conditional on the first-stage.

5.6 Continuation Values and Auction Entry

The same forward simulation yields each vessel’s continuation value (CV): the expected discounted value of continued operation accumulated as the log-sum of the daily choice problem over the horizon the vessel expects to remain active, discounted. I dollarize the continuation value with the structural revenue coefficient, $CV_i^{\$} = CV_i/\hat{\alpha}$, so that it is denominated in the same units as an auction bid.

The continuation value, not expected revenue, is the object a license-holder actually forfeits. A vessel deciding whether to sell weighs the entire discounted stream of option value from continued fishing, the ability to choose when and where to fish, and to stay ashore when conditions are poor, against the auction payment. Figure 7 traces the mean continuation value across active vessels over the sample, in structural utility units (panel A) and dollars (panel B), with the band showing the interquartile range. The continuation value is volatile from season to season, but shows no pronounced secular trend over the window. Its level and dispersion are descriptive context for the mechanism rather than identifying variation: the program’s effect on continuation values, and through them on participation, is recovered as the difference between the buyback and no-buyback arms of the counterfactual in Section 6, not from the observed path plotted here.




Figure 7: Mean continuation value across active vessels by season, in structural utility units (A) and dollars (B); the band is the interquartile range.

Table 7 shows the estimated link between continuation value and entry into the auction, with the sign the theory predicts: a higher continuation value lowers the probability that a vessel enters a round in order to sell ($\hat{\beta}_{CV} = -0.042$)¹².

¹²This sample is restricted to those who were active in a given year to minimize the probability of including a vessel that became “dormant”. In theory, if the auction worked exclusively and without leakage, so that the programmatic retirement of a license was a necessary and sufficient condition to exclude that license holder from this model, one could restrict the sample instead only to those who accepted a bid grant. But this is not the case: vessels go dormant for reasons outside of the auction, and some vessels appear to ticket landings in the round after they sell their license to TPWD. Nonetheless, a robustness check using this sample restriction instead is reported in Table A.4, and the magnitude is largely invariant to the selection of restriction.

Table 7: Auction entry and the continuation value

	(1) <i>CV</i>	(2) Expected revenue	(3) Both
<i>Variables</i>			
Continuation value CV_i	-0.0421* (0.0232)		-0.0419* (0.0239)
Expected revenue		-0.0045 (0.0108)	-0.0010 (0.0108)
Vessel length (ft)	0.0306 (0.0224)	0.0157 (0.0189)	0.0309 (0.0225)
<i>Fixed effects</i>			
Round	Yes	Yes	Yes
Primary bay	Yes	Yes	Yes
<i>Fit statistics</i>			
Observations	4,127	4,127	4,127
Pseudo R^2	0.042	0.036	0.042

Note:

Logit of the decision to enter (submit a bid) on the continuation value and/or expected revenue, with round and primary-bay fixed effects; standard errors clustered by vessel in parentheses. Column (1) enters CV alone, column (2) expected revenue alone, and column (3) both: CV is negative and stable across (1) and (3), while expected revenue is marginal and adds little once CV is included. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

6 Counterfactual Analysis: Spillovers, Offsets, and Targeting

6.1 Overview

I move toward a counterfactual analysis to quantify the participation distortion caused by stock externalities generated by the TPWD inshore shrimp buyback auction. Due to data limitations outlined in Section 3, harvest activity is only available at the bay-year level, not vessel-bay-day level, before 2008 when a large bulk of the program’s bids were placed. To construct counterfactual estimates at the program’s true scale required developing a unique approach of simulating a synthetic dataset of vessels disciplined by structural estimates of productivity, bay system choices, biological patterns, continuation values, and relationships with exogenous state variables in the 2008 to 2023 window, each using the structural estimates built in Section 5. I separately estimate the propensities of license-holders to make activity-related decisions to halt their activity outside of the auction, reactivate their license

after a period of dormancy, change their bay-system preferences in response to conditions and rival activity. This spillover mediation chain is illustrated in Figure 8, with Loop A showing how license retirement undermines the program’s recruitment and Loop B showing how the program reallocates persisting effort.


This section proceeds as follows. First, I describe the building of the synthetic fleet and the process of back-casting structural objects across the full 1997 to 2023 program, with 1996 included for pre-program calibration to an unobserved stock level, in Section 6.2. I then measure the participation distortion from a biological spillover in a forward simulation, comparing arms with and without the buyback, and decomposing compositional and selection effects of the participation margin in Section 6.3. Section 6.4 measures the structural offset — that is, the own- and cross-bay recapture of removed effort in response to license retirements. Sections 6.5 uses the offset estimate to analyze counterfactual policy recommendations for specified spatial targeting and design.

6.2 The Synthetic Fleet and Full-Horizon Back-Cast

The participation footprint isolated above is measured only over the window in which the active fleet is directly observed. Trip tickets matched to individual vessels begin in 2008; before that, from the program’s 1996 inception through 2007, the landings records are compiled only at the bay level, with no vessel identifiers to place a boat in the fleet. Yet the buyback’s activity is concentrated in exactly those early years (Figure 3), when removals were largest and, because more vessels were still fishing, the stock externality they generated was plausibly strongest. A counterfactual confined to 2008–2023 therefore observes the mechanism where its force has already ebbed.

To recover the program’s full-horizon footprint, I extend the structural objects estimated in Section 5—the production function, the location-choice probabilities, and the structural preferences—across the entire 1996–2023 horizon, filling the unobserved pre-2008 period by calibration to observed moments rather than by direct estimation. Because no matched panel reaches back to 1996, the fleet that fished the early rounds is reconstructed rather than read from data: I draw a synthetic 1996 fleet from the vessels that ever appear in the license and landings records, recover the initial stock by inverting the fleet’s simulated harvest onto observed 1996 landings, and propagate the result forward under the estimated dynamics, anchored at the far end to the observed program-end fleet. Algorithm B.1 in Appendix B details this construction step by step—the fleet draw, the joint inversion, and the anchor calibration; the paragraphs below give its logic and, for each object, state whether it is estimated, calibrated to a moment, or left unidentified and swept.

I recover unobservable stock conditions in 1996 by inversion. For a given bay, the observed 1996 harvest is a single number, but it reflects two unknowns: how many vessels fished and how healthy the stock was. One year’s harvest cannot separate them because more vessels and a healthier stock both raise landings. The active count is therefore pinned externally, from the license and retirement records, which just-identifies the stock: with the fleet fixed, I solve for the per-bay initial stock that makes the synthetic fleet’s simulated season harvest match observed 1996 landings. Because the five bays are coupled through the choice model—a healthier stock in one bay draws vessels toward it—the inversion solves the five bay-stocks jointly rather than one at a time. Because Gulf shrimp are short-lived, recruitment-driven



[†] Estimated but not identified in the baseline simulation: enters through vessel re-ranking, not the aggregate rate. The extensive-margin loop's magnitude is a bound, not a point estimate.

Figure 8: **Roadmap of the counterfactual mechanism.** Each round of license retirements (top) lowers cumulative depletion; survivors recapture 4–26% of the removed effort depending on local fleet density, the stock recovers within the season, and continuation values rise. Higher continuation values dampen three behavioural margins—fewer owners sell, fewer vessels go dormant, and dormant vessels reactivate, reallocating across bays through the choice model. **Loop A** (participation) undermines the program by slowing removals; **Loop B** (extensive margin) offsets the biological gain. Coefficients are structural estimates from Section 5; [†]-marked objects are estimated but not identified in the baseline simulation.

species, the model treats the stock as re-initialized each season and depleted within it; there is no cross-season biomass carryover, so removals propagate across years only through the reduced fleet. This matches the annual-crop character of the fishery; endogenizing multi-year biomass persistence would matter for longer-lived stocks and is left to future work.

Several inputs the in-window simulation reads do not exist for 1996. Buoy weather records do not reach back that far, so the season’s weather is seeded at a day-of-season climatological average; the unobservable productivity shock ζ is set to its estimated mean of zero; and the 1996 price is a bay-level annual revenue-per-pound proxy from the historical landings series. Each is a normalization absorbed into the recovered stock: the initial stock is net of the assumed weather, shock, and price scale. From the 1996 seed, the fleet evolves season by season under the estimated arrival and exit hazards, with the stock index propagating through the estimated depletion dynamics and carried across winters by an inter-season persistence parameter ρ . The level toward which stock reverts across seasons is calibrated to a second time point: the observed program-end (2018–2023) active fleet, by bay.

The production function, choice probabilities, structural preferences, and arrival and exit hazards are estimated (Sections 5). The 1996 stock and the reversion baseline are calibrated to observed time points. The object that remains to be swept is the share of the 1996 license pool that is latent rather than active, which the records do not determine and which I treat as a bounded assumption.¹³ Full mechanical detail is in Appendix B.

Because the early horizon is anchored to 1996 harvest and the late horizon to 2018–2023 fleet size, the recovered 1996 stock sits below the reversion baseline, so the simulated stock rises over the opening seasons as it moves from the first anchor to the second. This early rise is a seam between two calibration targets, not an ecological recovery, and it differences out the spillover, which is a between-arm comparison. I note that the fitted model reproduces the key observed moments: production and stock dynamics, location choice, and fleet decay; Appendix D documents the fit for each object.

A bidder retiring their licenses ostensibly means that their associated vessel is retired. But license holders may choose not to use their license when it’s available to them (serviced by the fact that about two-thirds of bidders submit zero harvest tickets in the same year), and new vessels appear regularly in the ticket data despite the 1995 moratorium on new license issuance. Two separate models are included to describe transitions of the observed fleet between dormancy and activity, detailed in Table 8. Natural exit is straightforward: vessels with licenses that have a lower associated continuation value are more likely to go dormant¹⁴. Reactivation is more challenging because dormant vessels are unobservable, thus do not have an estimated continuation value. For that reason I model it as a simple reduced-form hazard rate.

¹³The latent share affects the composition/backfill margins, but not the intensive-margin participation channel that is the paper’s headline, which is estimated among active vessels and is robust to the treatment of latent capacity.

¹⁴It is worth mentioning that while this model looks similar to the results in 7, the unit of observation is different, at vessel-round for the former and vessel-year for the later.

Table 8: Fleet-transition hazards (exit and reactivation)


	(1) Natural exit	(2) Reactivation
<i>Variables</i>		
Continuation value CV_i	-0.0244** (0.0124)	
Vessel length (ft)	-0.0105 (0.0107)	
Dormancy duration y_d		-0.8743*** (0.1347)
Dormancy duration ² y_d^2		0.0458*** (0.0131)
<i>Fixed effects</i>		
Year	Yes	Yes
Primary bay	Yes	Yes
<i>Fit statistics</i>		
Observations	3,152	1,150
Adj. pseudo R^2	0.007	0.142

Note: Reduced-form fleet-transition hazards used in the forward simulation, both logit with year and primary-bay fixed effects; standard errors clustered by vessel in parentheses. Column (1) is the probability that an active vessel exits for reasons other than the auction; Column (2) is the probability that a dormant vessel reactivates, as a function of years dormant y_d . * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

6.3 The Participation Distortion and Its Decomposition

The buyback’s success is itself a positive spillover onto the vessels that remain, operating through a four-link chain: (i) removals lower the cumulative harvest accumulating against the stock; (ii) lower cumulative harvest raises the stock index X_{jt} through the estimated mortality coefficient ϕ ; (iii) a higher stock raises each remaining vessel’s continuation value of holding a license; and (iv) higher expected revenue lowers the probability that a vessel enters an auction round. Links (i)–(iii) are governed by the structural production and stock dynamics of §5; link (iv) is the entry relation (Table 7).

I estimate the impacts of biological stock spillovers on auction participation using two forward simulations of the full program horizon that differ only in whether programmatic removals occur. Each arm’s fleet evolves endogenously—vessels arrive, fish, and exit under the estimated hazards, and the stock propagates through the estimated depletion—so the between-arm difference reflects spillover through an evolving fleet, the dynamic general-



Note: The between-arm participation difference decomposes additively into a within-vessel matched (behavioral) channel and a composition (selection) channel.

Figure 9: **Decomposition of the participation distortion.** *Top:* the between-arm difference in predicted auction entrants by year, split into the matched (behavioral) and composition (selection) channels. Bracketed interval is the 90% reduced-form confidence interval, propagating the (clustered) sampling uncertainty in β_{CV} through the participation recompute.

equilibrium form of the channel.

There are actually two separate channels that may distort participation in the two arms, which I refer to as a “matched” (behavioral) and “compositional” (selection) channel. The matched channel emerges when the same license in both arms has a different continuation value from the anticipated replenishment of stock health and reduces the license holder’s propensity to bid accordingly. The compositional channel emerges from what the program intended to do — remove marginal licenses and improve the aggregate quality of active vessels, q_i , who are then less likely to bid. I decompose the two channels in a similar fashion to an Oaxaca-Blinder decomposition [Oaxaca, 1973] by matching licenses in both arms and measuring differences in propensities to bid against licensees that appear in only the no-buyback arm.

The predicted participation distortion between the buyback- and no buyback-arms is illustrated in Figure 9. The appreciated stock replenishment from vessel removals increases the continuation value of remaining vessels by an average of 2.21 utils and decreased participation by about 7.6%. The participation gap is negligible in the opening rounds, widening substantially as retirements accumulate and the treated and counterfactual fleets pull apart, with the bulk of the effect concentrated in the program’s later, higher-volume years. This

gradual accrual is intrinsic: the absolute reduction in fishing capacity from removing these marginally-productive vessels reduces total harvest and improves stock health (see Figure C.1 for reference).

6.4 The Structural Offset: How Much Conservation a Removal Delivers

This chain is immediately suggestive of a lever that policymakers can employ to improve program efficiency — target bays (or patches) where the stock replenishment attracts less offset from those who remain. This section derives an identity for net capacity retirement contingent on endogenous patch density.

So far, I have described a demonstrable participation spillover in a program; retiring a vessel reduces capitalization in a bay, replenish stock health in that bay, license holders can harvest more shrimp at the same cost, increasing their continuation value and making exit less appealing. But retiring a vessel does not remove its effort one-for-one from the fishery. The resource the vessel leaves behind is partly reclaimed by the vessels that remain, which now fish a marginally richer ground and expand their own effort in response. The net conservation a buyback purchases is therefore the removed effort less this offset. This subsection derives the offset from the same primitives that drive the participation channel, and I show that it saturates in the size of the surviving fleet.

Within a season, a vessel’s harvest is proportional to its catchability and the local stock, $h_{ij} \propto q_i X_j$, and cumulative harvest depletes that stock, $\Delta \ln X_j = -\varphi \Delta H_j$, with the depletion elasticity governed by the Stage-1 cumulative-harvest coefficient $\beta_{\text{cum}} = 0.643$. Consider the marginal drop-one-vessel experiment underlying the counterfactual: remove an active vessel i from bay j and let the season re-solve. The removed vessel’s foregone harvest $\Delta = h_{ij}$ is a positive shock to the standing stock; the S_j vessels that remain each respond by fishing the richer ground more intensively, recapturing a share of Δ . Because the recapture is itself harvest, it re-depletes the stock, so each additional surviving boat recaptures a little less than the last. Aggregating the survivors’ responses gives a recapture that rises with, but saturates in, the surviving activity S_j :

$$\boxed{\text{offset}_j = \frac{\theta S_j}{1 + \theta S_j}, \quad \text{net capacity retired}_j = 1 - \text{offset}_j = \frac{1}{1 + \theta S_j}.} \quad (20)$$

The single parameter θ maps the structural depletion elasticity into the strength of the congestion-recapture feedback; estimated across bay-years it is $\theta \approx 0.345$, and the saturating form fits the simulated drop-one responses well (pseudo- $R^2 \approx 0.738$), materially better than a linear recapture, which overshoots in thick bays. Equation (20) carries the paper’s central tension in one line: where the surviving fleet is large, most of a removal is offset and net conservation is small; where the fleet is thin, a removal sticks. Figure 10 plots the simulated recapture against surviving activity by bay, with the fitted curve overlaid.

In principle the effort freed by a removal could be recaptured anywhere the removed vessel’s rivals fish, so the general object is a mobility matrix M whose (j, k) entry is the share of bay- j removals recaptured in bay k . Decomposing each removal’s recapture by where it lands based on own-bay (local congestion relief) versus cross-bay (spatial substitution) shows




Figure 10: Net capacity retired versus the size of the surviving fleet, one point per bay-year (70 bay-years, every other year 1997–2023). The solid curve represents the structural identity in (20)

that the cross-bay component is small throughout: own-bay recapture accounts for essentially all of the net offset in every bay (Table C.1). This is what the fishery’s loyalty structure would predict—vessels fish close to home, so the effort a removal frees is reclaimed mostly by boats in the same bay.

6.5 The Cost-Effectiveness of Removal Rules

A buyback’s job is to retire fishing effort, and the natural way to grade this effort is removals per dollar spent. One way to do this would be by crediting each retired license with its own effort—cost per license, or gross days removed. That accounting misses two things this paper has identified: The first is the dynamic offset of Section 6.4: the boats that remain recapture part of the freed stock, so net effort falls by less than the removed vessel’s own effort. The second is additionality: a share of the licenses retired belonged to vessels that were barely fishing, so buying them removed little real effort to begin with. Pricing both turns cost-effectiveness from an accounting number into a design problem with two administrable levers. I measure it as net trawl-days avoided per license retired, holding the budget fixed, and benchmark each rule against the naive per-license credit a static evaluation would assume (Figure 11).




Figure 11: Net trawl-days avoided per license retired, by removal rule, against the naive per-license credit assumed by static cost-per-license evaluation (dashed).

The program as run reduces trawl-days by 178 per retired license, only 21% of the naive credit of 830 days a static evaluation would assign. The gap is the two wedges combined; the offset piece alone described in Section 6.4, and the fact that roughly two-thirds of the licenses the program retired belong to vessels not fishing in the same season — non-additional purchases that removed little-to-no actual effort. Restricting eligibility to recently active vessels, a single administrable rule, raises net trawl-days reduced per license retirement to 301, or 36.3% of the naive credit, at the same license budget (the green bar in figure 11).

Because the offset is bay-contained and heterogeneous (Section 6.4), a license retired in a thin, low-offset bay removes more net effort than one retired in a thick, high-offset bay. Restricting removals to each bay in turn and re-running the program traces the spread in Figure 11: removing a license from San Antonio, a low-offset bay, reduces net trawl-days by 484 per license against 115 in Galveston, whose thick surviving fleet recaptures the largest proportion of removed effort¹⁵.

7 Conclusion

This paper structurally identifies a feedback through which the success of a buyback can weaken its own future participation: removals improve the resource, raise the continuation value of remaining agents, and make subsequent exit less attractive. Over the program’s 26-year history, TPWD received 2,515 bids; the counterfactual simulations imply that, absent the stock-replenishment spillover, roughly 204 additional bids would have been submitted. Importantly, these missing bids come from licenses associated with active fishing effort rather than from already-inactive capacity. Applying the program’s observed 44% retirement rate conditional on bidding implies that approximately 90 additional active licenses would have been retired in the absence of this feedback. This counterfactual is not merely hypothetical: a one-shot buyback would largely preclude the dynamic participation distortion by completing removals before beneficiaries can adjust their future participation to the improved resource.

The results also have direct implications for buyback design. The value of a retirement is heterogeneous: even before a purchase occurs, local resource density affects how much conservation that retirement ultimately delivers, because remaining users recapture part of the freed resource. Screening for recent activity is one obvious response to non-additional purchases, but such rules may themselves induce gaming—for example, an otherwise inactive license-holder could fish minimally to preserve eligibility. A more robust design principle is to target low-offset patches and, more generally, condition payments on expected net conservation rather than gross retirement. The same logic extends beyond fisheries wherever retirements alter the returns to remaining users: groundwater rights may be more valuable to retire where connected pumping responses are limited, while land-retirement programs may deliver more conservation where displaced production has fewer margins for reallocation.

Several limitations merit mention. First, vessel-level activity is unavailable before 2008, so the early program fleet is reconstructed rather than directly observed. The resulting back-cast is disciplined by observed harvest and late-period fleet moments, but inevitably relies more heavily on model structure than the 2008–2023 estimation window. Transitions between dormancy and active fishing are particularly difficult to recover because dormant vessels are, by definition, unobserved while inactive. Second, the estimated relationship between continuation value and auction participation is directionally consistent with the proposed mechanism but imprecisely estimated, with the baseline coefficient significant only at the 10% level. The counterfactual participation estimates should therefore be interpreted

¹⁵San Antonio’s ranking as the bay system with the most efficient per-license trawl-day reduction can seemingly contradict the fact that Matagorda and Aransas are less dense. Appendix Figure C.2 illustrates that while this is true, San Antonio also happened to have had the highest trawl-days per license initially, so while it is middle of the road at the intensive margin it dominated at the extensive margin.

together with their propagated uncertainty rather than as precisely identified point effects.

The counterfactuals in this paper evaluate the cost-effectiveness of alternative removal rules while holding the underlying auction mechanism fixed. They are therefore deliberately not welfare optima. A full welfare analysis would need to value conservation benefits against fisher surplus and allow both participation and bids to respond endogenously to changes in program rules. A related extension is to design transfers that internalize the externality a participant imposes on others. In the spirit of VCG mechanisms, such payments could in principle align private incentives more closely with the social value of retirement, although doing so may create substantial budget-balance and implementation challenges. Both exercises require a model of endogenous entry and strategic bidding and therefore constitute a distinct mechanism-design problem beyond the scope of the present paper. The structural objects recovered here provide much of the empirical foundation for that extension. For the present analysis, the contribution is narrower: to identify how much conservation alternative designs deliver per dollar once non-additional purchases and endogenous effort recapture are accounted for.

References

- Joshua K Abbott and James E Wilen. Regulation of fisheries bycatch with common-pool output quotas. *Journal of Environmental Economics and Management*, 57(2):195–204, 2009.
- Victor Aguirregabiria and Pedro Mira. Swapping the nested fixed point algorithm: A class of estimators for discrete markov decision models. *Econometrica*, 70(4):1519–1543, 2002.
- Victor Aguirregabiria and Pedro Mira. Sequential estimation of dynamic discrete games. *Econometrica*, 75(1):1–53, 2007.
- Rahel Aichele and Gabriel Felbermayr. Kyoto and carbon leakage: An empirical analysis of the carbon content of bilateral trade. *Review of Economics and Statistics*, 97(1):104–115, 2015.
- Peter Arcidiacono and Robert A Miller. Conditional choice probability estimation of dynamic discrete choice models with unobserved heterogeneity. *Econometrica*, 79(6):1823–1867, 2011.
- Karl M Aspelund and Anna Russo. Additionality and asymmetric information in environmental markets: Evidence from conservation auctions. *American Economic Review*, 116(8):2765–2805, 2026.
- Australian Government Department of Climate Change, Energy, the Environment and Water. Government water purchasing in the Murray–Darling basin, 2024. URL <https://www.dcceew.gov.au/water/policy/water-recovery/government-water-purchasing>.
- Patrick Bajari, C Lanier Benkard, and Jonathan Levin. Estimating dynamic models of imperfect competition. *Econometrica*, 75(5):1331–1370, 2007.
- Scott Barrett. Self-enforcing international environmental agreements. *Oxford Economic Papers*, 46(Supplement_1):878–894, 1994.
- Raphael Calel. Improving cost-efficiency of conservation auctions with joint bidding. *Journal of Environmental Economics and Policy*, 1(2):128–145, 2012.
- Pak-Sing Choi, Ana Espínola-Arredondo, and Félix Muñoz-García. Conservation procurement auctions with bidirectional externalities. *Journal of Environmental Economics and Management*, 92:559–579, 2018.
- Colin W Clark and Gordon R Munro. The problem of overcapacity. *Bulletin of Marine Science*, 70(2):473–483, 2002.
- Colin W Clark, Gordon R Munro, and Ussif Rashid Sumaila. Subsidies, buybacks, and sustainable fisheries. *Journal of Environmental Economics and Management*, 50(1):47–58, 2005.
- Sam Cunningham, Lori S Benneer, and Martin D Smith. Spillovers in regional fisheries management: do catch shares cause leakage? *Land Economics*, 92(2):344–362, 2016.

- Geret S DePiper. To bid or not to bid: The role of participation rates in conservation auction outcomes. *American Journal of Agricultural Economics*, 97(4):1157–1174, 2015.
- Godwin Kwabla Ekpe and Anna A Klis. Spillover effects in irrigated agriculture from the groundwater commons. *Environmental and Resource Economics*, 86(3):469–507, 2023.
- Ana Espinola-Arredondo. Green auctions: A biodiversity study of mechanism design with externalities. *Ecological Economics*, 67(2):175–183, 2008.
- Federal Reserve Bank of St. Louis. Federal reserve economic data (FRED), 2026. URL <https://fred.stlouisfed.org/>. Accessed: 2026-04-09.
- Paul J Ferraro. Asymmetric information and contract design for payments for environmental services. *Ecological Economics*, 65(4):810–821, 2008.
- Paul J Ferraro and Agnes Kiss. Direct payments to conserve biodiversity. *Science*, 298(5599):1718–1719, 2002.
- David A Fleming. Slippage effects of land-based policies: Evaluating the conservation reserve program using satellite imagery. *Papers in Regional Science*, 93:S167–S178, 2014.
- Allan Hernandez-Chanto and Andres Fioriti. Bidding securities in projects with negative externalities. *European Economic Review*, 118:14–36, 2019.
- Dan Holland, Eyjolfur Gudmundsson, and John Gates. Do fishing vessel buyback programs work: a survey of the evidence. *Marine Policy*, 23(1):47–69, 1999.
- V Joseph Hotz and Robert A Miller. Conditional choice probabilities and the estimation of dynamic models. *The Review of Economic Studies*, 60(3):497–529, 1993.
- Youxin Hu, John Kagel, Xiaoshu Xu, and Lixin Ye. Theoretical and experimental analysis of auctions with negative externalities. *Games and Economic Behavior*, 82:269–291, 2013.
- Ling Huang and Martin D Smith. The dynamic efficiency costs of common-pool resource exploitation. *American Economic Review*, 104(12):4071–4103, 2014.
- B Kelsey Jack, Carolyn Kousky, and Katharine RE Sims. Designing payments for ecosystem services: Lessons from previous experience with incentive-based mechanisms. *Proceedings of the National Academy of Sciences*, 105(28):9465–9470, 2008.
- Lewis C King and Jeroen CJM van den Bergh. Potential carbon leakage under the Paris Agreement. *Climatic Change*, 165(3):52, 2021.
- Joseph Kuehn. Estimating auctions with externalities: The case of USFS timber auctions. *Review of Economics and Statistics*, 101(5):791–807, 2019.
- Seokbeom Kwon. How does patent transfer affect innovation of firms? *Technological Forecasting and Social Change*, 154:119959, 2020.

- Douglas A Landis, Mary M Gardiner, Wopke van der Werf, and Scott M Swinton. Increasing corn for biofuel production reduces biocontrol services in agricultural landscapes. *Proceedings of the National Academy of Sciences*, 105(51):20552–20557, 2008.
- C. Liese, M. Emory, K. Furman, C. Gervasi, M. McPherson, S. Shoffler, and J. Walter. Snapshot: Economics of the gulf of america federal shrimp fishery. Technical Memorandum NOAA Tech. Memo. NMFS-SEFSC-795, National Oceanic and Atmospheric Administration, National Marine Fisheries Service, Southeast Fisheries Science Center, 2026. URL <https://www.fisheries.noaa.gov/s3//2026-03/Gulf-Shrimp-Economic-Snapshot-Report-Final-NMFS-SEFSC-795.pdf>.
- Charles F Mason and Andrew J Plantinga. The additionality problem with offsets: Optimal contracts for carbon sequestration in forests. *Journal of Environmental Economics and Management*, 66(1):1–14, 2013.
- Larry McKinney. TPWD to retain saltwater stamp surcharge. Press release, <https://tpwd.texas.gov/newsmedia/releases/?req=20040830a>, August 2004. Statement by TPWD Coastal Fisheries Division Director.
- NOAA Fisheries. American Fisheries Act pollock buyback program. <https://www.fisheries.noaa.gov/national/funding-financial-services/american-fisheries-act-pollock-buyback-program>, 7 2024a. Accessed: 2026-08-07.
- NOAA Fisheries. Fishing capacity reduction (buyback) programs. U.S. National Oceanic and Atmospheric Administration, 2024b. URL <https://www.fisheries.noaa.gov/national/funding-financial-services/fishing-capacity-reduction-programs>. Accessed: 2026-08-07.
- NOAA National Data Buoy Center. Historical standard meteorological data: Station 42035 (Galveston, TX). https://www.ndbc.noaa.gov/view_text_file.php?filename=42035h, 2026. Accessed: 2026-04-09.
- William Nordhaus. Climate clubs: Overcoming free-riding in international climate policy. *American Economic Review*, 105(4):1339–1370, 2015.
- Ronald Oaxaca. Male-female wage differentials in urban labor markets. *International Economic Review*, 14(3):693–709, 1973.
- Alexander Pfaff and Juan Robalino. Spillovers from conservation programs. *Annual Review of Resource Economics*, 9:299–315, 2017.
- Lisa Pfeiffer and C-Y Cynthia Lin. Groundwater pumping and spatial externalities in agriculture. *Journal of Environmental Economics and Management*, 64(1):16–30, 2012.
- Public Policy Institute of California. Managing water and farmland transitions in the San Joaquin Valley. Technical report, Public Policy Institute of California, 2023. URL <https://www.pplic.org/publication/managing-water-and-farmland-transitions-in-the-san-joaquin-valley/>.

- Michael J Roberts and Shawn Bucholtz. Slippage in the Conservation Reserve Program or spurious correlation? a comment. *American Journal of Agricultural Economics*, 87(1): 244–250, 2005.
- John Rolfe, Steve Schilizzi, Peter Boxall, Uwe Latacz-Lohmann, Sayed Iftekhar, Megan Star, and Patrick O’Connor. Identifying the causes of low participation rates in conservation tenders. *International Review of Environmental and Resource Economics*, 12(1):1–45, 2018.
- Andrew B Rosenberg and Bryan Pratt. Land use impacts of the Conservation Reserve Program: An analysis of rejected offers. *American Journal of Agricultural Economics*, 106(3):1217–1240, 2024.
- Mani Rouhi Rad, Dale T Manning, Jordan F Suter, and Christopher Goemans. Policy leakage or policy benefit? spatial spillovers from conservation policies in common property resources. *Journal of the Association of Environmental and Resource Economists*, 8(5): 923–953, 2021.
- Michael Elizabeth Sakas. To help refill two struggling underground aquifers, Colorado lawmakers are spending \$60 million to buy out water rights and farmland. CPR News, May 2022. URL <https://www.cpr.org/2022/05/13/colorado-aquifers-rio-grande-republican-river-basins-groundwater/>.
- James N Sanchirico and James E Wilen. Bioeconomics of spatial exploitation in a patchy environment. In *Fisheries Economics, Volume II*, pages 177–198. Routledge, 2020.
- Martin D. Smith. State dependence and heterogeneity in fishing location choice. *Journal of Environmental Economics and Management*, 50(2):319–340, 2005.
- Texas Parks and Wildlife Department. TPWD, CCA Texas team up to retire record number of inshore shrimping licenses. Texas Parks and Wildlife Department news release, Feb. 12, 2007. URL <https://tpwd.texas.gov/newsmedia/releases/?req=20070212a>.
- Texas Parks and Wildlife Department. Vessel/boat and outboard motor inquiry, 2026. URL <https://apps.tpwd.state.tx.us/tora>. Accessed: 2026-06-09.
- United States Coast Guard. Port state information exchange (PSIX), 2026. URL <https://cgmix.uscg.mil/psix/>. Accessed: 2026-06-09.
- U.S. Department of Agriculture, Farm Service Agency. Conservation reserve program (CRP) fact sheet. https://www.fsa.usda.gov/sites/default/files/2025-05/FSA_CRP_Factsheet_04_25e.pdf, 2025a. Accessed: 2026-02-02.
- U.S. Department of Agriculture, Farm Service Agency. Conservation reserve enhancement program (CREP), 2025b. URL <https://www.fsa.usda.gov/>.
- U.S. Department of Agriculture, Natural Resources Conservation Service. Agricultural conservation easement program: Wetland reserve easements, 2025. URL <https://www.nrcs.usda.gov/>.

U.S. House of Representatives. Voluntary grazing permit retirement act: Fact sheet, 2025. URL https://adamsmith.house.gov/sites/evo-subsites/adamsmith.house.gov/files/evo-media-document/vgpra-fact-sheet_2025.pdf.

JunJie Wu. Slippage effects of the Conservation Reserve Program. *American Journal of Agricultural Economics*, 82(4):979–992, 2000.

A Supplementary Estimation Results

A.1 Vessel Atomicity

Table A.1: Vessel concentration and the atomicity assumption

Bay system	Median vessel share (%)	99th pctl share (%)	Effective # vessels	Median # vessels	Top-vessel share (%)
Modeled bay systems					
Galveston Bay	0.52	4.8	55.9	118	4.3
Matagorda Bay	0.90	13.4	25.0	69	8.4
San Antonio Bay	1.70	11.5	20.8	40	8.7
Corpus Christi Bay	2.13	18.9	12.5	28	16.3
Aransas Bay	1.58	32.2	8.0	24	22.8
Other Texas bay systems (excluded from the model)					
Upper Laguna Madre	11.66	100.0	1.9	4	61.0
Lower Laguna Madre	56.89	100.0	1.4	2	85.2
Other	100.00	100.0	1.0	1	100.0
Sabine Lake	94.98	100.0	1.0	1	100.0
All Bays	0.88	29.5	13.2	30	15.0

Note: Each vessel’s share is its bay-shrimp landings in a bay-system-year divided by that bay-year’s total landings—its share of the *cumulative seasonal* harvest that determines the stock index, not of any single day. “Effective # vessels” is the inverse Herfindahl index ($1/\sum_i s_i^2$); the last three columns are medians across years within a bay. Bays are ordered densest to sparsest within each block. The upper block is the five modeled bay systems.

A.2 Stock-Index Components

Table A.2: Stock-index deterministic components

Term	Galveston	Aransas	Corpus Christi	Matagorda	San Antonio
Panel A. Within-season growth polynomial $G_j(t)$					
$\ln t$	0.701*** (0.073)	1.474*** (0.131)	0.373** (0.138)	0.996*** (0.106)	1.150*** (0.109)
$(\ln t)^2$	-0.707*** (0.049)	-1.073*** (0.065)	-0.606*** (0.067)	-0.898*** (0.057)	-0.969*** (0.059)
$(\ln t)^3$	0.237*** (0.013)	0.273*** (0.014)	0.219*** (0.014)	0.259*** (0.013)	0.262*** (0.014)
$(\ln t)^4$ (common)	-0.024*** (0.001)				
Panel B. Bay productivity					
$\ln A_j$ (bay FE)	-0.450	0.000	0.647	-0.049	0.321

Note: Deterministic components of the latent stock index (Equation (3)), from the Bartik-IV production regression (Table 3). Panel A is the within-season growth path $G_j(t)$: $\ln t$, $(\ln t)^2$, $(\ln t)^3$ vary by bay and $(\ln t)^4$ is pooled for power. Panel B is the bay productivity $\ln A_j$, absorbed as a bay fixed effect (no conventional standard error). The depletion coefficient is $\hat{\phi} = 0.643$ per million pounds (Table 3); the 16 year effects $\ln X_{0,\tau}$ (2008 the omitted base) range from -0.38 to 0.29 and are omitted for space. Standard errors in parentheses; * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

A.3 Bay-Shrimp Inverse Demand

A.4 Auction Entry: Risk-Set Robustness

Table A.3: Bay-shrimp inverse demand: local landings and the ex-vessel price

	(1) OLS	(2) IV
ln(bay landings)	−0.105 (0.012)	−0.092 (0.043)
Implied demand elasticity	−9.5	−10.9
First-stage F	—	20.9
Instruments (wind, wave height)	No	Yes
Bay fixed effects	Yes	Yes
Year-month fixed effects	Yes	Yes
Observations (bay-day)	9,693	9,693
Clusters (year-month)	106	106

Notes: The dependent variable is the log bay-shrimp ex-vessel price per pound; the regressor is log total bay landings at the bay-day level (bay food shrimp only). Column (2) instruments log landings with buoy wind speed and significant wave height, which shift the cost of fishing but are excluded from shrimp demand. Implied elasticity is the reciprocal of the coefficient. Standard errors clustered by year-month in parentheses.

Table A.4: Continuation value and auction entry: risk-set robustness

Dependent Variable: Model:	Bid	
	Activity-based (1)	Removal-censored (2)
<i>Variables</i>		
Continuation value CV_i	-0.0418* (0.0239)	-0.0442* (0.0242)
Vessel length (ft)	0.0309 (0.0225)	0.0308 (0.0225)
Expected revenue	-0.0010 (0.0108)	-0.0005 (0.0114)
<i>Fixed-effects</i>		
Round	Yes	Yes
Primary bay	Yes	Yes
<i>Fit statistics</i>		
Observations	4,127	3,905
Pseudo R ²	0.042	0.045

Clustered (vessel_key_std) standard-errors in parentheses

*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

Logit of the decision to bid on the continuation value, with round and primary-bay fixed effects; standard errors clustered by vessel in parentheses. Column (1) uses the risk set defined by observed fishing activity through the entry decision. Column (2) additionally censors vessel-rounds at or after a confirmed removal (a granted and accepted bid), dropping 249 vessel-rounds.

B Synthetic-Fleet Construction and Full-Horizon Back-Cast

This appendix gives the mechanical detail behind the back-cast summarized in Section 6.2: the 1996 fleet draw, the joint inversion that recovers the initial per-bay stock, the reversion-baseline calibration, and the season simulator and forward loop that the two policy arms share.

Algorithm B.1 Synthetic-fleet construction and full-horizon back-cast

Inputs: pool $\mathcal{P} = \{(q_i, \text{len}_i, b_i)\}$ (≥ 20 trips); active share σ_a ; size N ;
 estimated $\Theta = \{\phi, G_j, \theta, a_L, \lambda, \text{hazards}\}$;
 targets $H_{j,1996}^{\text{obs}}$ (1996 landings), F_j^{obs} (2018–2023 fleet).
Output: fleet $\mathcal{F}_{1996}$, stock $\mathbf{X}_{1996}$, reversion baseline $\{\bar{X}_j\}$.

Phase A. Draw the 1996 active fleet.

- 1: **for** each bay $j = 1, \dots, 5$ **do**
- 2: $s_j \leftarrow$ 1996 active share from historical landings, net of productivity
- 3: draw $n_j = \lfloor \sigma_a N s_j \rfloor$ vessels from $\mathcal{P}$, bay j , with replacement
- 4: **end for**
- 5: keep q_i ; impute length; inherit distances; assign loyalty a_L ; inherit revenue/lb

Phase B. Recover the 1996 stock by joint inversion.

- 6: seed $\bar{W}$ (climatology), $\zeta = 0$, price $\bar{p}_j$ $\triangleright$ normalizations absorbed into $\mathbf{X}_{1996}$
- 7: find $\mathbf{X}_{1996}$ solving $\text{SIMSEASON}(\mathcal{F}_{1996}, \mathbf{X}_{1996}; \bar{W}, 0, \bar{p}) = H_{\cdot,1996}^{\text{obs}}$
 5 coupled equations, 5 unknowns; $\mathcal{F}_{1996}$ fixed $\Rightarrow$ just-identified

Phase C. Calibrate the reversion baseline (late anchor).

- 8: **for** each bay j **do**
 - 9: set $\bar{X}_j$ so $\text{FORWARDSIM}(\text{no-removal})$ reproduces F_j^{obs}
 - 10: **end for**
 - 11: **return** $\mathcal{F}_{1996}$, $\mathbf{X}_{1996}$, $\{\bar{X}_j\}$
-

Algorithm B.2 Season simulator and forward back-cast (helpers for Algorithm B.1)

```

1: procedure SIMSEASON( $\mathcal{F}, \mathbf{X}_0; W, \zeta, p$ )
2:    $H_j^{\text{cum}} \leftarrow 0, H_j \leftarrow 0$  for all  $j$ 
3:   for each day  $t$  in the season do
4:     if  $t$  starts a sub-season then
5:        $H_j^{\text{cum}} \leftarrow 0$  ▷ depletion resets
6:     end if
7:      $X_{jt} \leftarrow G_j(t) - \phi H_j^{\text{cum}} + \dots$  ▷ compute  $X_{jt\_det}$ 
8:     each vessel picks a bay (or home) by the nested-logit CCP given  $X_{jt}, W, p$ 
9:     realize  $h_{ijt}$ ; add to  $H_j^{\text{cum}}$  and  $H_j$ 
10:  end for
11:  return  $\{H_j\}$  and realized activity
12: end procedure

13: procedure FORWARDSIM( $\text{arm}; \mathcal{F}_{1996}, \mathbf{X}_{1996}, \bar{X}, \Theta$ ; schedule  $R$ )
14:    $\mathcal{F} \leftarrow \mathcal{F}_{1996}, \mathbf{X} \leftarrow \mathbf{X}_{1996}$ 
15:   for  $\tau = 1996, \dots, 2023$  do
16:      $\{H_j\}, \text{activity} \leftarrow \text{SIMSEASON}(\mathcal{F}, \mathbf{X}; W_\tau, \zeta_\tau, p_\tau)$ 
17:     if  $\text{arm} = \text{treated}$  and  $\tau \in R$  then
18:       retire licenses per schedule, ranked by entry propensity
19:     end if
20:      $\mathcal{F} \leftarrow \text{CHURN}(\mathcal{F}, \text{activity}, \Theta)$  ▷ forward_churn
21:     record participation, stock,  $CV$  for  $\tau$ 
22:   end for
23:   return participation, stock, fleet,  $CV$  paths
24: end procedure

```




Figure C.1: Mean log-catchability of active vessels, by bay-year

C Counterfactual Supplementary Exhibits

Table C.1: Effort recapture by location, by bay.

Bay	Own-bay recapture	Cross-bay leakage	Net offset	Net capacity retired
Corpus Christi	3.8%	-0.1%	3.8%	96.2%
Aransas	6.2%	0.6%	6.7%	93.3%
San Antonio	8.6%	-1.3%	7.2%	92.8%
Matagorda	10.6%	2.7%	13.3%	86.7%
Galveston	26.5%	-0.1%	26.4%	73.6%

Note: Own-bay recapture (local congestion relief) and cross-bay leakage (spatial substitution) sum to the net offset; net capacity retired = 1–net offset. Bays ordered by net offset. Means over bay-years, 1997–2023




Figure C.2: Offset Decomposition

Notes: Shows the decomposition of what drives per-license conservation by bay. Net trawl-days avoided per license is the product of the additional (active) share of removals, the effort (days fished) of the retired vessels, and net capacity retained (1-offset).

D Model Validation

The counterfactual experiments of Section 6 rest on the estimated structural objects reproducing the behaviour observed in the data. This appendix reports one validation exercise for each layer of the model: the production function and stock dynamics (Figure D.1), the first-stage location-choice probabilities (Figure D.2), and the simulated fleet dynamics (Figure D.3).




Figure D.1: **Production-function fit at the trip and aggregate levels.** Observed against fitted log harvest; the dashed line is the 45° identity. *Panel A* is at the vessel-day (trip) level, where idiosyncratic trip noise is expected and the fit is correspondingly loose ($R^2 = 0.32$). *Panel B* aggregates to the bay-year level where the stock index and depletion dynamics operate and fits tight ($R^2 = 0.97$).




Figure D.2: **Calibration of the first-stage choice probabilities.** Reliability plot for the nested-logit conditional choice probabilities. Vessel-day location choices are grouped into bins by predicted probability; the plot shows the mean predicted probability against the observed choice frequency in each bin (marker area proportional to bin count). Points track the 45° line across the full range, indicating the estimated CCPs are well calibrated, suggesting predicted probabilities match realized choice frequencies at both rare and common alternatives.




Figure D.3: **Fleet decay: simulated versus observed active fleet.** *Panel A* overlays the simulated active-vessel path (line) on the observed active fleet (points) over the program horizon; the shaded region is the pre-2008 back-cast (calibrated, not estimated, Appendix B) and the dotted line marks Hurricane Harvey (2017). The simulation reproduces the decline closely through the mid-2010s, then over-retains vessels after Harvey, as the observed fleet continues to fall faster than the model. *Panel B* is the corresponding simulated-versus-observed scatter for the estimation window (2008–2023, $R^2 = 0.64$), showing the same post-Harvey flattening at the high end.